

AI humanistic reinterpretation of sharia

QUANTITATIVE TEXT ANALYSIS REPORT

DATE	September 12, 2026
ANALYSIS ID	01a095d6
SCOPE	Model: 5 groups
FEATURES	0 analytics applied

500	761	734.4K	0.21	26
DOCUMENTS	TOTAL CHUNKS	TOTAL WORDS	AVG. SENTIMENT	TOPICS

REPORT NOTICE

This report contains AI-generated text and analytical commentary. AI outputs may contain inaccuracies, fabricated references, or unintended bias. Content is provided for research and informational purposes only and should be independently verified before reliance.

AI-GENERATED SECTIONS

This report combines deterministic statistical analysis with narrative text generated by large language models (LLMs). The table below discloses every AI-generated section included in this report and the model that produced it. AI outputs may contain inaccuracies, fabricated references, or unintended bias and should be independently verified before reliance. All other sections are computed deterministically from the analyzed corpus.

AI-GENERATED SECTION	LOCATION IN REPORT	GENERATED BY
Executive Summary	Executive Summary section	gpt-5.6-luna (openai)
Composite Cross-Group Comparison	Composite Cross-Group Comparison section	gpt-5.6-luna (openai)
Conclusion & Recommendations	Conclusion and Recommendations section	gpt-5.6-luna (openai)
Deliverable	Deliverable section	gpt-5.6-luna (openai)

TABLE OF CONTENTS

- 1. Composite Cross-Group Comparison 5
- 2. Executive Summary 7
- 3. Deliverable 15
- 4. Corpus Overview 38
- 5. Conclusion and Recommendations 39

LIST OF TABLES

Table 1. Document composition by type in the current scope. 38

Table 2. Chunk length statistics in words for the current scope. 38

Table 3. Parameter coverage — Model: chunks per value in the current scope. 38

1. COMPOSITE CROSS-GROUP COMPARISON

Executive Overview

The five models converge on a humanistic reinterpretation of sharia centered on apostasy, religious conscience, minority protection, dignity, and the tension between classical rulings and contemporary rights. However, they differ substantially in how they organize that material. Claude Opus 4.7 is the most concentrated around legal pluralism and minority rights, GPT 5.5 is the most rights- and law-oriented, DeepSeek V4 Pro gives the greatest weight to formal legal foundations and punishment, Gemini 3.1 Pro preview foregrounds historical and interfaith framing, while Grok 4.3 emphasizes interpretive method, implementation, and scholarly pluralism.

The strongest common pattern is a movement away from purely punitive or state-centered discourse toward contextual interpretation. Across topic, lexical, phrase, co-occurrence, classification, and network results, terms such as *human rights*, *dignity*, *freedom*, *conscience*, *pluralism*, *justice*, and *protection* repeatedly connect legal questions to ethical and social consequences.

Key cross-group differences

- **Topic concentration varies sharply.** Claude assigns 73.49% of topic probability to "Islamic Legal Pluralism and Minority Rights," leaving only 14.02% for apostasy and 12.49% for blasphemy. By contrast, DeepSeek distributes attention more evenly, with 49.67% on dignity and pluralism, 26.07% on punishments and human rights, and 24.25% on apostasy. GPT 5.5 is more differentiated: minority rights account for 43.13%, state institutions and rights for 23.37%, and women's rights for 17.33%. Gemini places 37.85% on human dignity in interfaith relations and 28.79% on apostasy, whereas Grok spreads attention across blasphemy and minority protections (31.16%), humanistic scholarship (28.97%), and gender justice topics.
- **Lexical salience distinguishes legal intensity from ethical framing.** Apostasy is the leading TF-IDF term for Claude (5.439), DeepSeek (9.902), GPT (11.507), and Grok (6.494), while Gemini makes *law* dominant (8.820), followed by apostasy (7.438). DeepSeek's word frequency also emphasizes legal formalism: *law* appears 486 times and *apostasy* 421 times. GPT is the most expansive corpus, with *law* at 1,169 and *ethical* at 1,029, while Grok's leading terms are *interpretation* (350) and *justice* (326), indicating a more methodological orientation.
- **Sentiment is uniformly positive but not equally so.** Grok is the most consistently affirmative, with 98.7% positive and no negative cases. Gemini follows at 91.0% positive and 2.6% negative. GPT records 86.5% positive and only 4.5% negative, while Claude and DeepSeek are more mixed: 77.0% and 77.9% positive, with negative shares of 13.0% and 16.9%, respectively. These differences suggest that punishment, apostasy, and coercion generate more tension in Claude and DeepSeek than in the other outputs.

Cross-feature patterns

Several findings reinforce one another. GPT's classification assigns 61.5% of chunks to Foundations of Islamic Law, while its frequent phrases emphasize *sharia* (1,138), *justice* (556), *freedom* (318), and *conscience* (282). This combines a strong legal base with an explicitly ethical vocabulary. DeepSeek similarly combines Foundations of Islamic Law (35.1%) with Religious Rights (31.8%), and its co-occurrences connect *property* with protection of life, intellect, and lineage at 0.5383 and *dignity* with *karāmah* at 0.3259.

Gemini and Grok show the clearest plurality-oriented signals. Gemini's classification gives Schools of Law and Theology 16.0%, far above GPT's 1.5% and Claude's 1.0%; its keyword matches include Hanafi, Maliki, and Hanbali in 83.3%, 83.3%, and 82.7% of chunks. Grok's strongest co-occurrences include *individual right-pluralism* at 0.3593 and *majority-minority* at 0.3185, while its n-grams include *interfaith relation* at 0.858 and *jurisprudence of minorities* at 0.689.

A notable tension concerns coherence versus breadth. Claude's LDA is highly concentrated but reports low-coherence and overlapping topics, whereas Grok and Gemini offer broader thematic differentiation but also show multiple overlapping or low-coherence topics. These results caution against treating topic labels as discrete doctrines; the models repeatedly blend law, ethics, history, and rights discourse.

Similarities and convergences

All models produce dense, integrated conceptual networks: every network has one connected component containing all 50 nodes. Density ranges from 0.8522 for Claude to 0.9780 for GPT, and clustering ranges from 0.8719 to 0.9798. Thus, differences concern emphasis rather than isolated vocabularies.

Across n-grams and co-occurrences, *human right*, *human dignity*, *historical context*, *universal declaration*, *freedom of conscience*, and *maqāṣid* recur. Chunking results likewise place *sharia*, *human rights*, apostasy, freedom, conscience, dignity, and justice among the most frequent phrases for nearly every group.

Strategic implications

Stakeholders should treat model outputs as complementary interpretive profiles rather than competing final authorities. GPT is useful for broad legal-ethical synthesis; DeepSeek for explicit doctrinal and punishment analysis; Claude for minority protection and pluralism; Gemini for madhhab-aware historical comparison; and Grok for implementation, interpretive method, and interfaith application.

A robust workflow should combine these profiles, preserve juristic disagreement, distinguish doctrine from contemporary state practice, and require human review whenever outputs address apostasy penalties, blasphemy, gender justice, or minority status. The quantitative convergence around dignity, conscience, protection, and contextual interpretation offers a strong foundation for a karāmah-centered and maqāṣid-oriented discourse, provided that positive sentiment is not mistaken for juridical validity or ethical adequacy.

2. EXECUTIVE SUMMARY

2.1. Claude Opus 4.7 Analysis

Executive overview

The corpus contains 500 documents, 761 chunks, and 734,447 words, divided across five model groups with 100 documents each. The groups are unequal in chunk and word volume: GPT 5.5 contributes 200 chunks and 214,396 words, while Claude Opus 4.7 contributes 100 chunks and 108,608 words. This makes raw cross-model counts descriptive rather than directly comparable; proportions, group-specific keyness, and normalized rates would provide a stronger basis for claims about model behavior.

Across the available evidence, the models construct sharia primarily through a human-rights and reform-oriented frame rather than through a narrowly punitive frame. LDA assigns 73.49% of the reported topic weight to Islamic legal pluralism and minority rights, compared with 14.02% for apostasy laws and religious conscience and 12.49% for blasphemy offenses and punishments. Classification points in the same direction: Reform and Legal ethics accounts for 43% of the aggregate classification table, Foundations of Islamic Law 24%, and Religious rights 20%, while Gender and Women accounts for only 10%. This suggests that gender, dignity, interfaith relations, and juristic reasoning are present but are often subordinated to a broader debate about pluralism, conscience, and reform.

Five higher-order themes organize the corpus. First, pluralism and minority protection connect religious freedom, citizenship, apostasy, and interfaith coexistence. Second, maqāṣid-oriented ethical reasoning links harm prevention, public welfare, and protection of life, intellect, lineage, and property. Third, juristic and theological plurality brings Sunni, Shia, Maliki, Hanafi, Shafi'i, Hanbali, Ja'fari, Ash'ari, Maturidi, and Mu'tazili references into one interpretive field. Fourth, punitive law and conscience form a contested boundary around apostasy, blasphemy, repentance, compulsion, and state authority. Fifth, gender, dignity, and humanistic reform appear as a comparatively smaller but normatively important cluster.

The strongest cross-feature relationship is between thematic dominance, phrase structure, and conceptual association. The phrase counts for sharia (361), human rights (202), apostasy (139), fiqh (110), freedom (108), belief (85), conscience (75), women (64), minorities (42), justice (40), dignity (20), and Karāmah (14) show that the discourse repeatedly places legal concepts beside rights and conscience language. Co-occurrences such as dīn_compulsion-non_compulsion (NPMI 0.5718), maslaha-public_welfare (0.5447), darar-harm (0.6212), and property-protection_life_intellect_lineage (0.5575) indicate that ethical reinterpretation is being linguistically organized around non-coercion, welfare, harm reduction, and protected human interests.

Model differences are clearest in sentiment and classification framing. Gemini 3.1 Pro preview has the highest average sentiment at 0.2877, closely followed by Grok 4.3 at 0.2854; Claude Opus 4.7 is lowest at 0.1085. GPT 5.5 contributes the largest volume and is more strongly associated with Foundations of Islamic Law, with 123 classified chunks, whereas DeepSeek V4 Pro has 49 Religious rights chunks and Gemini has 25 Schools of Law & Theology chunks. These patterns suggest different rhetorical emphases, but they do not establish that one model is more humane or more legally accurate: VADER sentiment is a rough lexical signal, and classification categories impose the analytic frame.

The corpus over-represents pluralism, rights, apostasy, and reform discourse, while gender-specific analysis is comparatively thin. Women appears 64 times, but women's rights and gender justice appear only 10 and 14 times

respectively in the phrase inventory; interfaith relations are represented through terms such as Christians (20), Jews (8), coexistence (17), and Jews-Christians-Zoroastrians co-occurrence (NPMI 0.4862), but the evidence does not show comparable depth for lived experiences of religious minorities. Karāmah appears 14 times, suggesting that dignity-centered language is present but not yet a dominant organizing concept. Future analysis should compare equal-sized chunks, report per-model normalized frequencies and keyness, inspect the missing TF-IDF and word-frequency outputs, validate classification performance, and manually evaluate whether apparent humanistic language reflects substantive pluralism or repeated safety-oriented wording.

2.2. DeepSeek V4 Pro Analysis

Executive overview

The DeepSeek V4 Pro segment contains 100 documents, 154 analyzed chunks, and 144,485 words. Across these materials, sharia is represented primarily through a human-rights and religious-liberty debate rather than through a balanced treatment of all five requested issues. Topic modeling assigns 49.67% prevalence to "Sharia, Dignity, and Pluralism," 26.07% to "Sharia Punishments and Human Rights," and 24.25% to "Apostasy and Religious Liberty." This distribution suggests that dignity, protection, pluralism, and minority rights provide the segment's main interpretive frame, while apostasy, punishment, and blasphemy supply its most contested legal cases.

Higher-order thematic synthesis

Five macro-themes organize the evidence. First, **dignity, protection, and pluralism** connect "dignity," "karāmah," "minority," "pluralism," and "protection"; the strongest topic-level distinctive terms include dignity at 2.5963 and karāmah at 2.1998. Second, **apostasy, conscience, and coercion** centers on apostasy, punishment, repentance, freedom, and the state, with apostasy especially distinctive at 9.712. Third, **juristic plurality and interpretive method** is visible in the recurring presence of Hanafi, Maliki, Shafi'i, and Hanbali references, as well as maqāṣid, ijtihad, consensus, Mu'tazila, and Ash'ari theology. Fourth, **gender and family justice** appears through woman, women's rights, divorce, testimony, marriage, and child marriage, but it is less prominent than religious-rights discourse. Fifth, **state authority and the public regulation of belief** links punishment, blasphemy, execution, law, political authority, and religious freedom.

The five humanistic issues are therefore unevenly represented. Religious rights and interpretive reasoning are strongly foregrounded: the classification output assigns 31.8% of chunks to Religious rights and 35.1% to Foundations of Islamic Law, while keyword coverage reaches 72.1% for human rights, 69.5% for freedom, and 70.1% for jurisprudence. Human dignity is conceptually central, but interfaith relations are more marginal and usually appear through minority protection, dhimmi/dhimma terminology, Medina, citizenship, and pluralism rather than through sustained analysis of relationships between specific faith communities. Gender and women's rights are also comparatively underdeveloped: the Gender and Women class contains 5 chunks, or 3.2%, although woman appears in distinctive topic vocabulary and marriage occurs in 46 chunks, or 29.9%.

Cross-feature interpretation

The topic, co-occurrence, classification, and phrase results point to a consistent interpretive pattern: DeepSeek V4 Pro tends to mediate sharia as a negotiable ethical-juridical tradition whose legitimacy is tested against dignity, protection, conscience, and pluralism. The co-occurrence pair property-protection_life_intellect_lineage has the highest reported NPMI, 0.5383, while maqāṣid-protection_life_intellect_lineage reaches 0.3486 and dignity-

karāmah reaches 0.3259. These associations suggest that the corpus connects legal reasoning to protection-oriented objectives and human worth, rather than presenting rules as detached technical commands.

The same pattern appears in phrase packaging. "Sharia" occurs 683 times, "apostasy" 274 times, "human rights" 126 times, "women" 111 times, "freedom" 101 times, and "conscience" 62 times. The repeated phrases "the death penalty" at 79 occurrences, "religious freedom" at 20, "gender justice" at 16, "minority rights" at 18, and "human dignity" at 29 show how the model frames legal questions through recognizable contemporary rights vocabulary. However, these counts measure recurrence, not argumentative quality or endorsement.

Debates, consensus, and gaps

The dominant debate concerns whether classical rulings on apostasy, blasphemy, punishment, and gender can be reinterpreted through maqāṣid, historical context, conscience, and universal rights. The strongest topic overlap is between Topic 3 and Topic 6 at 0.5789, sharing dignity, scholar, right, minority, authority, human, justice, state, and woman. Additional overlaps of 0.55 and 0.5217 connect Topic 3 with religious protection and with law, history, jurisprudence, and human rights. These overlaps indicate a shared argumentative vocabulary rather than sharply separate themes; they also explain why the model's topics are difficult to distinguish conceptually.

A provisional consensus is visible around protection, justice, pluralism, and the need to distinguish historical legal contexts from contemporary state enforcement. Co-occurrences such as the four Sunni schools with "major_sunni_school" at 0.4884, fiqh al-aqalliyyāt with minority at 0.4179, and the Universal Declaration of Human Rights with "human_rights_declaration" at 0.4306 suggest that the corpus treats legal plurality and modern rights discourse as connected resources. Yet the data do not establish agreement among jurists or sources; they show repeated linguistic association in AI-mediated texts.

The principal gaps are sustained interfaith analysis, detailed gender jurisprudence, and differentiated treatment of madhhab disagreement. Interfaith language is present, including interfaith_relation, Christians, Jews, dhimmi, and Medina, but it does not form a dominant topic. Gender terms appear in topic vocabularies and co-occurrences such as man-woman at 0.3844 and divorce-woman at 0.3111, but the low Gender and Women classification share indicates that gender justice is often subordinated to the broader rights-and-law frame. The segment also provides no model-to-model comparison because the scoped analysis excludes the global baseline and contains only DeepSeek V4 Pro.

Methodological cautions and recommendations

The LDA model selected seven topics, but two topics were flagged as low-coherence: Topic 1 at 0.2769 and Topic 3 at 0.2786. Topic 3 has 49.67% prevalence while combining punishment, blasphemy, dignity, reform, protection, and pluralism, which suggests that its apparent dominance may partly reflect a broad or internally mixed topic. Several topic pairs also overlap at 0.50 or above, so the five macro-themes are more analytically defensible than treating all seven topics as independent domains.

The very dense text network, with 50 nodes, 1,055 edges, density 0.8612, average clustering 0.8888, and one connected component, indicates strong lexical integration. It may represent genuine conceptual cohesion, but it may also reflect repeated AI-generated framing and the shared five-token co-occurrence window. Similarly, positive sentiment covers 77.9% of chunks, but VADER detects lexical and rhetorical tone, not juristic validity; positive labels should not be read as approval of particular punishments or interpretations.

For deeper research, the segment should be reanalyzed with a smaller set of theory-driven themes, separate source-level and model-generated material, and explicit coding for the five humanistic issues. Comparative evaluation across the other models should test whether DeepSeek V4 Pro is unusually rights-oriented, unusually repetitive around apostasy, or simply responding to a shared prompt structure. Manual review should also examine whether references to Hanafi, Maliki, Shafi'i, Hanbali, Mu'tazila, and Ash'ari traditions represent substantive disagreement or merely lists of authorities.

2.3. GPT 5.5 Analysis

Executive overview

The GPT 5.5 segment contains 100 documents, 200 chunks, and 214,396 words. Its discourse is strongly organized around the relationship between Islamic law, religious freedom, and humanistic ethical reform: topic modeling assigns 43.13% of topic weight to religious minority rights, 23.37% to state institutions and Islamic rights, and 17.33% to women's rights in Islamic law. Text classification similarly places 61.5% of chunks in Foundations of Islamic Law, 19.0% in Reform and Legal ethics, and 15.0% in Religious rights. Taken together, these results suggest that GPT 5.5 frames humanistic reinterpretation primarily as an internally reasoned juridical and ethical problem, rather than as a narrowly gender-centered or interfaith-relations discussion.

Higher-order thematic clusters

Four overlapping macro-themes organize the segment. First, **juridical foundations and interpretive plurality** connect sharia, fiqh, jurisprudence, the Qur'an, revelation, consensus, and the Hanafi, Maliki, Shafi'i, and Hanbali traditions. Second, **religious liberty and minority protection** links apostasy, conscience, coercion, religious freedom, dhimma, equal citizenship, and state law. Third, **humanistic ethical reform** brings together dignity or karāmah, justice, mercy, public welfare, accountability, pluralism, and human rights. Fourth, **gender and family justice** appears through women, family law, marriage, inheritance, divorce, and gender justice, but it is less prominent in the thematic and classification distributions than the other clusters. Interfaith relations form a fifth, smaller cluster, visible through Christians, Jews, interfaith relations, Muslim-non-Muslim relations, covenant, and cooperation.

The strongest cross-method pattern is the coupling of legal reasoning with moral concepts. Co-occurrence scores link modern human rights with the Universal Declaration of Human Rights (NPMI 0.3916; count 48), dignity with karāmah (0.3240; count 48), justice with mercy (0.2838; count 45), and human with moral responsibility (0.2581; count 44). The association of coercion with conscience (0.2634; count 66) and belief with coercion (0.2508; count 52) indicates that religious rights are constructed through a recurring debate about the limits of state or communal power, not simply through abstract praise of tolerance.

The segment's positive sentiment classification should not be read as proof of endorsement. VADER labels 86.5% of chunks positive, 9.0% neutral, and 4.5% negative; in an academic and reform-oriented corpus, positive lexical cues may reflect terms such as justice, protection, dignity, mercy, and reform rather than emotional approval. The very dense text network--density 0.9780, average clustering 0.9798, one connected component containing all 50 nodes--shows a highly integrated vocabulary, but it may also indicate saturation from repeated framing and shared terminology. Because no other model group is included, claims of uniqueness are limited to the internal profile of GPT 5.5; direct comparative superiority or difference cannot be established.

Implications, gaps, and recommendations

Observed evidence suggests that GPT 5.5 offers a relatively conciliatory, pluralism-oriented representation of sharia, in which juristic reasoning mediates between inherited legal categories and contemporary human-rights language. The discourse gives substantial attention to apostasy, blasphemy, punishment, conscience, and state authority, but the relatively low representation of Gender and Women in classification (3.0%) and the lower phrase count for women (125) compared with sharia (1,138), justice (556), and freedom (318) indicate that gender justice is not the principal organizing frame. Interfaith relations are also comparatively marginal: interfaith relations occur 21 times, Muslim-non-Muslim relations 19 times, and interfaith cooperation 12 times in the extracted phrase inventory.

The main research gaps are therefore not total absences but uneven development. Women's rights appear through family law and marriage, yet the available outputs do not show comparable depth on women's interpretive authority, lived experience, economic agency, or intra-Muslim diversity. Interfaith discourse is present but thin relative to minority-rights and apostasy debates. The corpus also provides no direct evidence about Shi'i, Ibadi, female, non-clerical, or non-Muslim source voices; these should be treated as provenance questions requiring source-level review rather than inferred absences from vocabulary alone.

For further analysis, researchers should compare GPT 5.5 with the other model segments using identical keyness, topic-stability, and classification-quality outputs. They should also inspect document-level provenance and quoted material, test whether repeated prompt templates produced lexical convergence, and separate discussions of descriptive classical law from normative reform proposals. A multilabel or human-coded evaluation focused on the five core issues would improve on the single-label taxonomy, while targeted sampling of passages containing women, interfaith relations, *karāmah*, apostasy, and juristic disagreement would reveal whether the model preserves genuine *ikhtilāf* or merely repeats pluralist vocabulary.

2.4. Gemini 3.1 Pro preview Analysis

Executive overview

This segment contains 100 documents, 156 chunks, and 139,434 words focused on humanistic reinterpretations of sharia. The corpus presents a strongly rights-oriented and reform-sensitive discourse: the largest classification groups are Reform and Legal ethics at 49 chunks (31.4%) and Foundations of Islamic Law at 43 (27.6%), followed by Religious rights at 30 (19.2%), Schools of Law and Theology at 25 (16.0%), and Gender and Women at 9 (5.8%). These figures suggest that gender is addressed within a broader legal-ethical framework rather than functioning as the dominant organizing category.

Higher-order thematic synthesis

Five overlapping macro-themes organize the segment. **Humanistic legal reform and maqāṣid reasoning** links humanistic interpretation, family-law reform, justice, mercy, dignity, and the objective of law. **Religious freedom and coercion** brings together apostasy, blasphemy, conscience, punishment, repentance, and the distinction between religious belief and political treason. **Pluralism and interfaith citizenship** connects minority rights, the Ottoman millet system, dhimmī history, international human rights, and coexistence. **Juristic plurality and theological method** is represented by Hanafi, Maliki, Shafi'i, and Hanbali references alongside *kalām*, Ash'ari, Maturidi, and Mu'tazili terms. **Gendered justice and agency** appears through women, property, female autonomy, gender justice, and family-law reform, but has less classificatory coverage than the other themes.

The LDA output reinforces these clusters but also shows redundancy. "Apostasy Laws and Religious Freedom" is repeated across several topic rows and has the highest displayed topic value, 28.79, while "Human Dignity in Interfaith Relations" has 37.85 and "Islamic Pluralism and Interfaith Relations" has 12.65. The repeated apostasy topic indicates a central debate, but it may also reflect multiple near-duplicate formulations rather than six fully distinct thematic structures. Because the model used seven topics, and the supplied output does not include stability, coherence, or perplexity values, the topic count should be treated as a useful organizing device rather than definitive evidence of seven independent themes.

Cross-feature interpretation

The strongest conceptual relationships are not simply between rights terms; they connect legal reasoning to ethical restraint. Co-occurrence is especially strong for *maqāṣid al-sharī'a* and objective law (NPMI 0.5887, count 54), justice and mercy (0.5571, count 23), offender and repent (0.5550, count 31), and capital punishment and repent (0.4786, count 21). This pattern suggests that the segment frames difficult rulings through purposes, mercy, repentance, and contextual judgment rather than through punishment alone. The association of apostasy with death penalty (0.3244, count 21), treason (0.3079, count 24), and classical ruling (0.2814, count 22) shows that the corpus also preserves the juridical and political terms of the contested debate.

The phrase inventory gives the discourse a clear lexical center: "sharia" occurs 1,135 times, "Islamic law" 378 times, "human rights" 168 times, "freedom" 150 times, "apostasy" 125 times, "women" 81 times, and "human dignity" 73 times. The high repetition of "sharia," together with 1,079 edges among 50 network nodes, density of 0.8808, average clustering of 0.9024, and one connected component, indicates a highly integrated vocabulary. It may reflect conceptual cohesion, but it may also signal prompt-driven repetition or formulaic AI-generated exposition. These results therefore support claims about recurring framing, not necessarily about the independent importance of each concept.

The 91.0% positive sentiment rate, compared with 6.4% neutral and 2.6% negative, should not be read as evidence that the corpus endorses every reform position. VADER classifies sentence-level lexical tone, and academic writing about dignity, freedom, justice, and protection naturally contains positive terms. The result more plausibly indicates a constructive, normative register; quoted positions, descriptions of historical violence, and unresolved disagreements may not be adequately distinguished from authorial evaluation.

Representation, gaps, and recommendations

The segment strongly represents Sunni legal-school plurality: Hanafi and Maliki each occur in 130 chunks, Hanbali in 129, and Shafi'i in 125; Ash'ari appears in 100, Maturidi in 98, and Mu'tazila in 94. It also repeatedly connects classical jurisprudence to universal human rights, including international human rights and the Universal Declaration of Human Rights (NPMI 0.3411, count 31). This is a significant humanistic framing, but the available outputs do not show whether non-Sunni, Shi'i, Ibadi, women-authored, or geographically diverse sources receive comparable substantive treatment. Their absence from the reported prominent vocabulary should be treated as a research gap, not proof of complete exclusion.

Gender and women's rights are comparatively underrepresented in the classification output at 5.8%, despite meaningful terms such as "gender justice" (50 occurrences), "women's rights" (23), "female scholars" (14), "female autonomy" (7), property (7), and "gender equity" (11). This suggests that gender is present but often subordinated to general rights, reform, or jurisprudential framing. Interfaith relations and dignity are more visible through "interfaith relations" (43), "pluralism" (39), "minorities" (26), "minority rights" (25), and "human dignity" (73), yet the

data do not establish whether lived experiences of religious minorities are represented or merely discussed as legal categories.

For further analysis, the corpus should be tested with source-level provenance metadata, author and school identification, and separate comparisons of gender, religious freedom, interfaith relations, dignity, and juristic reasoning. Topic stability and classification macro-F1 should be reported before making strong claims about thematic boundaries. A manually coded sample should also distinguish descriptive historical discussion, normative legal argument, and AI-generated synthesis, while expanding the taxonomy to include marginalized schools, Shi'i and Ibadi perspectives, queer and disability-related dignity questions, and women's own interpretive agency.

2.5. Grok 4.3 Analysis

Executive overview

The Grok 4.3 segment contains 100 documents, 151 analytical chunks, and 127,524 words. Its discourse is strongly organized around interpretive jurisprudence, religious-minority protection, and ethical reform rather than around gender or interfaith relations as independent themes. The clearest evidence is the topic structure: "Comparative Juristic Interpretation" links jurisprudence, pluralism, scholars, protection, women, and rights, while "Religious Minority Rights" repeatedly combines apostasy, belief, punishment, coercion, conscience, and historical context.

Across methods, the segment presents sharia as a negotiated ethical-legal field. Co-occurrence links "human rights declaration" with "universal declaration human rights" at NPMI 0.4959, "individual right" with "pluralism" at 0.3593, and "dignity" with "karāmah" at 0.2903. These associations suggest that Grok 4.3 frames humanistic reinterpretation through a bridge between Islamic jurisprudential concepts and contemporary rights language, while retaining internal legal vocabulary such as the Hanafi, Maliki, Hanbali, and Shafi'i schools.

Five higher-order macro-themes consolidate the overlapping outputs. First, **juristic plurality and interpretive authority** joins comparative interpretation, the four Sunni legal schools, consensus, historical reasoning, and scholarly authority. Second, **freedom of conscience and minority protection** joins apostasy, blasphemy, coercion, religious freedom, and state power. Third, **dignity-centered ethical reform** connects karāmah, justice, mercy, welfare, equity, and human rights. Fourth, **gender justice as a marginally developed application** appears through women, inheritance, equality, and gender justice, but its classification presence is very low. Fifth, **community, governance, and coexistence** links community, state, Medina, interfaith relations, pluralism, and protection.

The classification results reinforce this framing: Foundations of Islamic Law accounts for 58 chunks (39.2%), Reform and Legal Ethics for 52 (35.1%), Religious Rights for 24 (16.2%), Schools of Law and Theology for 12 (8.1%), and Gender and Women for only 2 (1.4%). This distribution should not be read as a direct measure of intellectual importance, because the predefined taxonomy and keyword rules shape the result. It does, however, indicate that the segment gives much more space to legal foundations and reform ethics than to sustained gender analysis.

The network is highly cohesive: 50 nodes form one connected component with 1,171 edges, density 0.9559, and average clustering 0.9588. This cohesion supports the interpretation that rights, jurisprudence, dignity, community, and historical reasoning are repeatedly discussed together. It may also indicate saturation or repeated framing rather than a genuinely diverse conceptual structure. The chunking results show especially strong repetition of "sharia" (717 noun-phrase occurrences), followed by "justice" (136), "scholars" (122), "belief" (100),

"communities" (96), "apostasy" (76), and "dignity" (60), suggesting a discourse built around named concepts and entities more than explicit action-oriented argumentation.

Sentiment should be interpreted cautiously. VADER labels 149 of 151 chunks as positive (98.7%), 2 as neutral (1.3%), and none as negative. In an academic and normative corpus, this likely reflects the positive vocabulary of justice, protection, dignity, freedom, and reform rather than demonstrable approval of particular legal positions. It should not be used as evidence that the model resolves contested issues harmoniously.

The main research gap is not the absence of rights language but the uneven development of the five requested issues. Religious rights and juristic interpretation are richly represented; human dignity is conceptually connected to *karāmah* and justice; interfaith relations appear but are less prominent; and gender and women's rights are substantially underrepresented as an analytical category. The segment also offers limited evidence about Shi'i, Ibadi, non-Sunni, Muslim feminist, lived-experience, or non-Muslim voices. These absences cannot be confirmed as literal source omissions from the supplied summaries, but the low gender-classification share and strong Sunni-school keyword coverage warrant targeted verification.

Actionable interpretation should therefore combine the strong jurisprudential architecture with deliberate balancing. Future prompts or sampling should require parallel treatment of women's rights, interfaith relations, and human dignity alongside apostasy and minority rights; distinguish descriptive presentation of classical rulings from normative endorsement; and request explicit comparison across madhhabs, theological schools, modern reformist approaches, and affected communities. A second analytical run should report per-document topic proportions, TF-IDF terms, classifier validation metrics, and source provenance so that repeated AI-generated formulations are not mistaken for independent scholarly convergence.

3. DELIVERABLE

3.1. Claude Opus 4.7

Purpose: Academic Literature Review Synthesizer

Corpus overview

The corpus contains 100 resources represented by 100 chunks and 108,608 words for Claude Opus 4.7. The analysis used stopword removal and lemmatization, with topic modeling, TF-IDF, word frequency, sentiment, n-gram, co-occurrence, classification, network, and phrase-extraction methods.

The corpus is strongly oriented toward legal-ethical discussion. The most frequent terms include "ruling" (350 occurrences), "historical" (336), "law" (289), "ethical" (265), "jurisprudence" (257), "apostasy" (246), "protection" (237), "jurist" (234), "scholar" (222), and "pluralism" (217). These frequencies indicate recurring attention to legal interpretation, historical contextualization, rights protection, and disagreement among jurists, but frequency alone does not establish normative priority.

Research themes

The quantitative evidence supports one broad, interconnected field of inquiry organized around rulings, communities, principles, justice, protection, and practice. The synthesis labels this broad concentration "Ruling / Community / Principle," with distinctive vocabulary including "woman," "punishment," "ethical," "dignity," "scholar," "reform," and "pluralism." However, the LDA export is internally inconsistent: it reports a configured topic count of seven, displays duplicated topic identifiers and labels, and lists 26 topics in the coverage summary.

Accordingly, the topic output can support lexical orientation but not reliable claims about distinct topic prevalence. The reported dominant-topic distribution assigns 75% of chunks to topic 0, 14% to topic 1, and 11% to topic 2, while other displayed prevalence values and repeated records conflict with this distribution. No valid coherence, perplexity, seed-stability, or C_NPMI and U_Mass results are available, and the analysis itself flags low coherence and overlapping topics.

Gender and women's rights

Gender is a visible but comparatively smaller concentration within the corpus. "Woman" occurs 169 times and is among the top TF-IDF terms, while gender-related topic vocabulary includes "justice," "right," "reform," "authority," "dignity," "property," "inheritance," "patriarchal," and "historical context."

The lexical pattern suggests that Claude frames women's rights through the relationship between jurisprudential authority, historical interpretation, protection, and reform. This is an interpretation of the vocabulary rather than evidence that the model consistently endorses a particular gender jurisprudence. The classification taxonomy assigns 10 chunks to "Gender and Women," and all 10 receive positive VADER labels, but this small category should not be treated as a validated estimate of corpus-wide gender emphasis.

Religious rights

Religious rights constitute the clearest substantive concentration. "Apostasy" is the highest TF-IDF term at 6.6745 and occurs 246 times, while the apostasy-related lexical field includes "belief," "punishment," "coercion,"

"conscience," "state," "repent," "political," and "religious freedom." A separate blasphemy and punishment vocabulary includes "offense," "sabb al-nabi," "insult," "ḥudūd," "intent," and "repent."

The evidence distinguishes several related but non-identical problems rather than treating religious rights as a single issue. Apostasy is connected to freedom of belief, conscience, dissent, and state enforcement; blasphemy is connected to offense, speech, punishment, intent, and prophetic honor; and coercion is connected to the possibility of non-compulsion. The classification assigns 20 chunks to "Religious rights," although classifier performance and confidence measures are unavailable.

Sentiment varies within this issue. The apostasy-related topic group has an average compound score of 0.0014, with seven positive, five negative, and two neutral chunks. The blasphemy and punishment group is negative, averaging -0.1092, with eight negative and three neutral chunks. This contrast indicates that the discourse becomes more evaluatively charged around punishment and blasphemy than around apostasy as a broader question of belief and legal status.

Interfaith relationships

Interfaith relations appear through historical institutions, minority protection, and contemporary rights discourse. Relevant vocabulary includes "minority" (192 occurrences), "religious," "dhimma," "dhimmi," "Jews," "Christians," "Zoroastrians," "millet," "covenant," "coexistence," and "jurisprudence of minorities." Resource-level distinctive terms include "dhimma_contract," "variation," "al-Andalus," "medieval," "confessional," "Christian," "Malaysia," and "Almohad."

The n-grams "jurisprudence of minorities" (57 occurrences), "religious minorities" (52), "minority and majority" (57), and "Muslim-majority state" (78) indicate sustained attention to the legal and political positioning of non-Muslim communities. The corpus also contains resources explicitly comparing identity-based and dignity-based approaches to interfaith relations and examining historical protections under dhimmi arrangements. These signals support an interpretation of interfaith discourse as negotiated between inherited status-based frameworks, historical cases of coexistence, and dignity- and rights-oriented models.

Co-occurrence patterns reinforce this interpretation but should be understood as proximity rather than demonstrated conceptual equivalence. "Jews_Christians" co-occurs with "Zoroastrians" with an NPMI of 0.4862 across 16 instances, while "fiqh_al_aqalliyāt_jurisprudence" co-occurs with "minority" with an NPMI of 0.4427 across 33 instances. "Cairo Declaration Human Rights" and "Universal Declaration Human Rights" co-occur 34 times with an NPMI of 0.4025, linking Islamic and international rights vocabularies in the corpus.

Human dignity

Human dignity functions as a cross-cutting ethical vocabulary rather than an isolated subject. "Dignity" occurs 144 times and appears alongside "human," "right," "protection," "justice," "minority," "woman," "ethical," and "pluralism." The phrase "human dignity" occurs 76 times, while "protection of life" occurs 55 times.

The network places "protection," "ethical," "minority," "pluralism," and "right" among the most connected or central terms, alongside generic legal terms such as "law," "ruling," and "historical." This pattern suggests that Claude's discourse frequently mediates legal questions through ethical and protective language. Centrality does not, however, demonstrate that dignity is the model's normative foundation; frequent and generic terms can become network hubs because of their broad applicability.

The co-occurrence of "maslaha" and "public_welfare" has an NPMI of 0.5447 across 21 instances, while "darar" and "harm" have an NPMI of 0.6212 across 15 instances. These pairings indicate that welfare, harm prevention, and protection provide important conceptual resources for humanistic reinterpretation, although contextual concordances would be needed to determine how consistently they are used in argument.

Interpretive juristic reasoning

Interpretive reasoning is represented through "jurisprudence" (257 occurrences), "interpretation" (205), "jurist" (234), "scholar" (222), "historical," "context," "verse," "quranic," "uṣūl al-fiqh," "maqāṣid," and "public welfare." The n-grams "historical context" (89), "uṣūl al-fiqh" (55), "Quran foundation" (64), and "sunni juristic" (62) indicate a recurring effort to connect textual sources with legal method and historical circumstances.

The co-occurrence structure displays explicit intra-tradition and cross-tradition linkages. "Ash'ari_Maturidi_school" and "Mu'tazila" co-occur 16 times with an NPMI of 0.5153; "Ja'fari" and "Shī'ī" co-occur 18 times with an NPMI of 0.4875; and "Shia" and "Sunni" co-occur 28 times with an NPMI of 0.4785. "Malik_ibn_Anas" and "Maliki_school" co-occur 15 times with an NPMI of 0.5063, while "Hanbali" and "Ibn_Taymiyya" co-occur 20 times with an NPMI of 0.4216.

These patterns support an interpretation of Claude's discourse as engaging juristic plurality through schools, theologians, historical context, and methodological concepts. They do not establish equal treatment of all madhhabs or demonstrate the validity of the interpretations attributed to them. The model-specific keyness profile is especially important: relative to the comparison corpus, Claude overuses "reformist," "sharī'a," "read," "reading," "majority_state," "maliki_school," "jurist," "verse," and "critique."

At the same time, Claude underuses "justice," "law," "human_rights_declaration," "adapt," "community," "respect," "interpretation," "individual," "quranic_foundations_sharia," and "public_welfare" relative to the rest of the corpus. These results describe relative framing, not absolute absence. The comparison group, its balance, and multiple-comparison adjustments are not fully documented, so the keyness findings should be treated as directional evidence of model-specific emphasis.

Key concepts from TF-IDF and phrase analysis

TF-IDF foregrounds "apostasy," "ruling," "ethical," "law," "woman," "jurisprudence," "protection," "scholar," "jurist," "pluralism," "religious," "punishment," "right," "minority," "theological," "sharī'a," "interpretation," "dignity," and "reform." The combination shows that Claude's representation of sharia is not limited to doctrine or state law; it repeatedly connects legal rulings with ethical evaluation, historical scholarship, minority protection, gender, and reform.

The most frequent phrases include "human right" (522), "Universal Declaration" (115), "declaration of human" (114), "humanistic reading" (102), "Cairo Declaration" (95), "right in Islām" (93), "historical context" (89), "Muslim-majority state" (78), and "freedom of conscience" (46). These phrases indicate a strong rights-based and comparative framing of sharia discourse.

The n-gram field requires methodological caution. Several values reported as "significance" exceed the stated theoretical NPMI range of -1 to +1, including 1.1255, 1.2576, and 1.3369. They should not be interpreted as valid NPMI scores until the output is relabeled or recomputed with NPMI and G^2 reported separately.

Author and institution mapping

The available analytics do not provide a reliable author, institution, organization, or entity-extraction map. The entity landscape fields for people, regions, and organizations are empty, so no evidence-based ranking of scholars or institutions can be supplied.

The vocabulary does identify intellectual and jurisprudential references, including Joseph Schacht, Ignaz Goldziher, Ibn Taymiyya, Mālik ibn Anas, Sunni, Shia, Ja'fari, Hanbali, Maliki, Ash'ari-Maturidi, and Mu'tazila. These are lexical references in the corpus, not evidence of author influence, institutional affiliation, citation counts, or balanced representation.

Paper clustering

The classification taxonomy assigns the 100 chunks to six categories: 43 to "Reform and Legal ethics," 24 to "Foundations of Islamic Law," 20 to "Religious rights," 10 to "Gender and Women," two to "State, Politics & Security," and one to "Schools of Law & Theology." These assignments provide a practical organizational scheme for the literature review, but no cross-validated macro-F1, class-level precision and recall, confusion matrix, confidence estimate, or class-balance assessment is available.

The LDA output cannot currently substantiate a stable clustering of papers. Topic identifiers and labels are duplicated, prevalence values are inconsistent, overlap pairs include contradictory pairings, and the macro-theme summary reports 100% prevalence for a broad seven-topic cluster despite the separate dominant-topic distribution. The most defensible clustering interpretation is therefore taxonomy-based and lexical: reform and legal ethics form the largest assigned group, followed by foundations of law and religious rights, while gender and women's rights receive narrower explicit coverage.

Conceptual architecture

The co-occurrence and network results indicate a densely connected vocabulary centered on historical interpretation, rulings, law, ethics, jurisprudence, scholars, apostasy, pluralism, and interpretation. The network contains 50 nodes and 1,044 edges, with density of 0.8522 and average clustering of 0.8719. Because the network uses a five-token window and includes frequent generic terms, these density measures should not be interpreted as proof of a fully integrated conceptual system.

The strongest network hubs and bridges include "historical," "protection," "ruling," "scholar," "minority," "ethical," "law," "right," "jurist," and "jurisprudence." This architecture is consistent with a model that commonly moves from historical context and juristic authority toward questions of protection, rights, and ethical evaluation. It does not by itself establish that these concepts are causally or normatively prioritized.

Sentiment profile

Across all 100 chunks, the mean VADER compound score is 0.1085, with 77 positive, 10 neutral, and 13 negative chunks. The overall positive balance should not be read as evidence that the discourse offers a positive ethical evaluation of sharia. VADER is lexicon-based, and terms such as "protection," "rights," and "justice" may contribute positive scores even when they occur in discussions of coercion, punishment, or controversy.

The topic-aligned sentiment differences are more informative than the aggregate average. The broad Topic 1 grouping has an average of 0.1605, with 70 positive and five neutral chunks; the apostasy-related Topic 2 grouping averages 0.0014; and the blasphemy and punishment Topic 3 grouping averages -0.1092. These results suggest a

shift from constructive or reform-oriented language toward neutrality and negativity as the discourse addresses religious penalties and offenses.

Research gaps

The corpus provides weaker or less reliable coverage in several areas.

- Gender and women's rights receive identifiable attention, but the 10-chunk classification category is substantially smaller than reform, foundations, and religious-rights categories.
- State, politics, and security are represented by only two chunks, making claims about governance, enforcement, and political authority unstable.
- Schools of law and theology are assigned only one chunk by the classifier, despite substantial lexical evidence of Maliki, Hanafi, Hanbali, Shia, Ja'fari, Ash'ari-Maturidi, and Mu'tazila references.
- Human dignity is distributed across many themes rather than measured as a distinct category, limiting direct assessment of its comparative prominence.
- Interfaith relations are visible through minority, dhimmi, covenant, coexistence, and historical case vocabulary, but no validated issue-level prevalence or balanced regional comparison is available.
- The corpus does not provide source balance, author balance, resource length normalization, regional metadata, or author and institutional mapping.
- No direct comparison with the other models is available because the run includes only Claude Opus 4.7 and disables the global baseline.

The most important technical gap is the need to repair or rerun the topic model with unique topic identifiers, a single documented topic count, valid prevalence values summing to one, representative passages, and complete coherence, perplexity, and seed-stability reporting. The n-gram output also requires correction, and the classifier requires performance and uncertainty reporting before its categories can be treated as empirical prevalence.

Synthesis

Claude Opus 4.7 represents sharia primarily as a field of interpretive and ethical negotiation rather than as a single fixed legal code. Its vocabulary repeatedly connects rulings with historical context, juristic authority, pluralism, minority protection, human rights, dignity, reform, and the tension between individual conscience and state enforcement.

The five requested humanistic issues are interrelated but not equally visible. Religious rights are the clearest concentration, organized around apostasy, blasphemy, punishment, coercion, conscience, and state authority. Interfaith relations are mediated through historical protection, minority jurisprudence, dhimmi and dhimma concepts, and comparisons between identity-based and dignity-based models. Human dignity operates as a cross-cutting ethical vocabulary linked to protection, rights, women, minorities, and public welfare. Gender and women's rights appear through justice, authority, property, inheritance, protection, and reform. Interpretive juristic reasoning is expressed through madhhab and theological plurality, uṣūl al-fiqh, maqāṣid, historical context, and competing readings of scripture and precedent.

The model-specific profile adds an important qualification. Relative to the comparison corpus, Claude foregrounds reformist language, the term sharī'a, reading and rereading, majority-state contexts, Maliki references, jurists, verses, and critique. It relatively underuses justice, adaptation, community, respect, individual, public welfare, and human-rights-declaration language. This suggests a discourse that emphasizes interpretive method and reformist engagement, while not necessarily foregrounding the broader social and welfare vocabulary used by other models.

Overall, the evidence supports describing Claude's output as humanistic in its recurring attention to dignity, protection, pluralism, historical context, and interpretive disagreement. That conclusion should remain provisional: malformed LDA records, unavailable validation metrics, uncertain corpus comparability, lexicon-based sentiment scoring, and the absence of cross-model alignment limit claims about topic stability, normative consistency, or comparative superiority.

3.2. DeepSeek V4 Pro

Purpose: Academic Literature Review Synthesizer

Corpus overview

The DeepSeek V4 Pro segment contains 154 text chunks drawn from 100 resources and approximately 144,485 words. The analysis applied topic modeling, TF-IDF, word-frequency and keyness analysis, sentiment analysis, n-gram and co-occurrence analysis, text classification, network analysis, and phrase extraction.

The corpus is strongly organized around legal and interpretive vocabulary. The most frequent terms include *law* (598 occurrences), *apostasy* (477), *ruling* (432), *right* (376), *historical* (327), *state* (323), *Quran* (311), *scholar* (283), *ethical* (281), *protection* (276), *jurisprudence* (275), and *interpretation* (265). These frequencies describe the segment's vocabulary, but they should not be treated as uniquely DeepSeek-specific emphasis without comparison with the other models.

Research themes

The available topic output indicates three broad semantic formations rather than a reliably validated seven-topic solution. The dominant assignment covers 79 of 154 chunks (51.3%), a second covers 40 chunks (26.0%), and a third covers 35 chunks (22.7%). However, the export contains repeated topic identifiers, inconsistent topic counts, and incompatible keyword sets assigned to the same identifiers; the LDA quality metrics and stability score are also absent. These results therefore support thematic orientation, not definitive topic prevalence.

Across the corpus, DeepSeek's responses connect legal rulings to historical context, scriptural interpretation, state authority, minority protection, and human rights discourse. A recurring interpretive field links *law*, *ruling*, *jurisprudence*, *scholar*, *Quran*, *historical*, *right*, and *protection*. A second field is more concentrated around apostasy, punishment, coercion, conscience, blasphemy, and religious freedom. A third brings together dignity, pluralism, minorities, women, community, and maqāṣid-oriented protection.

Gender and women's rights

Gender appears as a substantive but comparatively limited classification category. Only 5 of the 154 chunks were assigned to "Gender and Women," so conclusions based on that classifier are descriptive and highly unstable. Nevertheless, *woman* occurs 216 times, appears among the highest-weighted TF-IDF terms, and is associated with *justice*, *reform*, *jurisprudence*, *authority*, *dignity*, and *gender* in topic vocabulary.

The co-occurrence pair *man-woman* appears 30 times with an NPMI of 0.3844, indicating a recurring relational framing of gender. Resource-level signals identify *patriarchal readings*, *gender*, *Qur'an*, *divorce*, *inheritance*, *property*, and *gender equality*. The corpus therefore represents women's rights primarily through reinterpretation of patriarchal readings, protection, family-law questions, and maqāṣid reasoning, although the limited classified coverage prevents a strong claim about the consistency of this orientation.

Religious rights

Religious rights are the most clearly represented of the five requested issues in the classification output, with 49 chunks assigned to "Religious rights." The issue is also the strongest distinctive lexical formation: *apostasy* is the leading TF-IDF term and occurs 477 times, while *religious freedom*, *freedom of conscience*, *coercion*, *blasphemy*, *punishment*, *state*, and *belief* recur throughout the corpus.

DeepSeek's comparative keyness profile shows overrepresentation of *apostasy* relative to the rest of the corpus, with a log-likelihood statistic of 51.85 and a log-ratio of 0.62. It also overuses *hadith* (log-ratio 1.68), *verse* (1.30), *Quran* (0.63), and *reinterpretation* (1.32). This combination suggests that religious-rights questions are frequently approached through scriptural, hadith-based, historical, and juristic argument rather than through abstract rights language alone.

The apostasy-related topic has the lowest reported mean sentiment score, 0.0141, with 17 negative, 15 positive, and 3 neutral chunks. This should be understood as rhetorical polarity rather than an ethical judgment: VADER sentiment cannot establish whether the responses endorse religious freedom, coercive punishment, or a particular theological position.

Interfaith relationships

Interfaith relations are not reported as a separate classifier category, and no reliable chunk count is available. They are nevertheless visible through the recurring vocabulary of *minority*, *dhimmi*, *community*, *pluralism*, *religious*, *protection*, and *diversity*. The term *dhimmi* is overrepresented relative to the rest of the corpus, with a log-ratio of 1.11, while the co-occurrence of *fiqh al-aqalliyyāt jurisprudence* and *minority* has an NPMI of 0.4179 across 22 observed pair instances.

A resource-level signal explicitly centers "karāmah" in everyday interfaith interactions, alongside *dignity*, *honor*, and spiritual sharing. This indicates that interfaith relations are framed through minority jurisprudence, protection, dignity, and practical interaction. The evidence does not establish how consistently DeepSeek distinguishes interfaith dialogue from broader religious-rights or minority-protection discourse.

Human dignity

Human dignity is a central normative vocabulary item, occurring 232 times and appearing among the highest TF-IDF terms. It is closely connected to *right*, *protection*, *justice*, *minority*, *woman*, and *pluralism*. The pair *dignity-karāmah* occurs 29 times with an NPMI of 0.3259, providing direct evidence that the Arabic-Islamic concept of *karāmah* is used as a semantic bridge between humanistic language and Islamic ethical discourse.

The corpus also connects *maqāṣid* reasoning with protection objectives. The pair *maqāṣid-protection_life_intellect_lineage* occurs 18 times with an NPMI of 0.3486, while *property-protection_life_intellect_lineage* is the strongest reported association, with 38 instances and an NPMI of 0.5383. These patterns support an interpretation of dignity as articulated through protection, welfare, and the objectives of law, although the relative underrepresentation of *ethical* and *justice* in keyness analysis qualifies any claim that DeepSeek foregrounds broad ethical language more than peer models.

Interpretive juristic reasoning

Interpretive reasoning is embedded throughout the segment's legal vocabulary. The most frequent terms include *ruling*, *jurisprudence*, *interpretation*, *scholar*, *historical*, *Quran*, *verse*, *hadith*, *maqāṣid*, and *ijtihad*. DeepSeek's keyness profile particularly emphasizes *hadith*, *maqāṣid*, *verse*, *reinterpretation*, *major Sunni school*, *Medina*, and *Shafi'i*. This indicates a relatively strong scriptural and school-based register in the visible comparison with the rest of the corpus.

The co-occurrence structure also explicitly names juristic and theological plurality: *Hanafi*, *Maliki*, *Shafi'i*, and *Hanbali* are grouped with "major Sunni school," while *Mu'tazila* is associated with *rationalist* and *reason*, and *Ash'ari* with *theological school*. These associations demonstrate that the corpus recognizes multiple interpretive traditions. However, the available data do not compare how each school is represented, whether disagreements are treated symmetrically, or whether one methodology is privileged.

A significant qualification is that *interpretation* is underrepresented relative to the rest of the corpus, with a log-ratio of -0.74, and *disagreement* is markedly underrepresented, with a log-ratio of -3.04. DeepSeek therefore uses many juristic and scriptural markers while comparatively less often foregrounding explicit disagreement, broad interpretation, or interpretive contestation.

Key concepts from TF-IDF and keyness

The highest TF-IDF terms are *apostasy*, *law*, *ruling*, *right*, *dignity*, *historical*, *state*, *Quran*, *ethical*, *scholar*, *protection*, *jurisprudence*, *interpretation*, *pluralism*, *human*, *woman*, and *punishment*. This vocabulary positions the segment at the intersection of Islamic legal reasoning, human-rights discourse, state regulation, and historical interpretation.

The strongest DeepSeek-specific contrasts are more informative than raw frequency. Relatively overrepresented terms include:

- *hadith* with a log-ratio of 1.68.
- *reinterpretation* with a log-ratio of 1.32.
- *verse* with a log-ratio of 1.30.
- *maqāṣid* with a log-ratio of 1.20.
- *dhimmi* with a log-ratio of 1.11.
- *major Sunni school* with a log-ratio of 1.09.

Relatively underrepresented terms include *ethical* with a log-ratio of -1.01, *justice* at -1.16, *protection* at -0.71, *interpretation* at -0.74, and *disagreement* at -3.04. Taken together, the profile suggests a discourse that is relatively more scriptural, historical, and juristic, while using less explicit language of broad ethical framing, disagreement, and generalized protection than the comparison corpus.

Paper clustering and network structure

The resource-level evidence supports several overlapping clusters: gender and patriarchal interpretation; apostasy, blasphemy, and religious expression; hudud and punishment; religious freedom and conscience; karāmah and interfaith interaction; and maqāṣid-oriented gender and human-rights reasoning. These clusters should be treated as interpretive groupings rather than validated paper clusters because the topic export does not provide a canonical document-by-topic matrix or reliable topic-quality statistics.

The text network contains 50 selected nodes and 1,055 edges, with density of 0.8612 and average clustering of 0.8888. The most prominent weighted and PageRank nodes are *law*, *ruling*, *apostasy*, *right*, *Quran*, *historical*, *state*, *interpretation*, *protection*, and *jurisprudence*. This shows that legal, scriptural, historical, and state-related terms connect many of the five humanistic issues.

The very high density should not be interpreted as evidence that all concepts are equally central. The graph is limited to 50 selected terms and uses a sliding-window construction, which can mechanically increase connectivity. It is best read as evidence of an interconnected legal-humanistic vocabulary rather than as proof of clearly separated conceptual communities.

Author and institution mapping

No reliable author, institution, regional, organizational, or people-level entity mapping is available in the analytics output. The synthesis therefore cannot identify leading scholars, institutions, or geographic traditions in the resource set. The presence of terms such as *Hanafi*, *Maliki*, *Shafi'i*, *Hanbali*, *Mu'tazila*, and *Ash'ari* supports analysis of intellectual traditions, but not attribution of claims to specific authors or institutions.

Research gaps

The most evident coverage gap is the absence of a validated five-issue comparison. Religious rights have 49 classified chunks, whereas gender and women's rights have only 5; no equivalent coverage counts are available for interfaith relations, human dignity, or interpretive juristic reasoning. This imbalance makes it difficult to determine whether apparent emphases reflect substantive model behavior, taxonomy design, or the composition of the resources.

Additional gaps concern interpretive plurality and ethical balance. The corpus names multiple schools and traditions but does not compare their treatment, evidences underuse of *disagreement*, and does not provide representative excerpts for evaluating epistemic humility or the preservation of *ikhtilāf*. Resource-level and chunk-level analyses also lack a resource-weighted robustness check, so repeated chunks from the same resource may influence frequency, sentiment, co-occurrence, and keyness results.

The topic model requires technical validation before it can support stronger clustering claims. The configured topic count is 7, but the output exposes 15 entries, repeated identifiers, and a reported coverage count of 26 topics; coherence, perplexity, and stability values are missing. The n-gram output also requires audit because values labeled as significance include values above the stated NPMI upper bound of 1, making those scores unsuitable for substantive interpretation until NPMI and G^2 are separately reported.

Synthesis

DeepSeek V4 Pro represents sharia-related humanistic reinterpretation through a highly interconnected legal and scriptural vocabulary. Its responses repeatedly connect law and ruling to apostasy, state authority, historical context, Quranic and hadith evidence, juristic schools, minority status, dignity, and maqāṣid. Among the five requested issues, religious rights are the most visibly developed in the classification data, while gender is present through women's rights, patriarchal readings, family law, and gender justice but has limited classified coverage.

The segment's distinctive profile is not simply "humanistic" in the abstract. Relative to the rest of the corpus, it more strongly emphasizes hadith, verses, maqāṣid, reinterpretation, dhimmi status, and Sunni-school terminology, while using less ethical, justice, protection, interpretation, and disagreement vocabulary. The most defensible interpretation is therefore that DeepSeek constructs humanistic reinterpretation primarily through

historically situated scriptural and juristic reasoning, with dignity and karāmah providing an important ethical bridge.

The evidence supports a plural-aware discourse, but not yet a fully demonstrated account of juristic plurality. Multiple schools and rationalist and theological traditions appear in the vocabulary, while explicit disagreement and comparative school analysis remain underdeveloped. Stronger conclusions about DeepSeek's ideological orientation, comparative humanism, or relative support for gender equality and religious freedom require a harmonized cross-model pipeline, validated topic assignments, resource-level weighting, classifier diagnostics, and a dedicated five-issue dashboard.

3.3. GPT 5.5

Purpose: Academic Literature Review Synthesizer

Corpus overview

The GPT 5.5 segment contains 200 text chunks drawn from 100 resources and 214,396 words. The analysis uses lemmatized, stopword-filtered text and covers topic modeling, TF-IDF, word frequency, sentiment, n-grams, co-occurrence, classification, network analysis, and phrase chunking.

The corpus is classified into five categories: Foundations of Islamic Law accounts for 123 chunks, Reform and Legal Ethics for 38, Religious Rights for 30, Gender and Women for 6, and Schools of Law and Theology for 3. These counts describe the available analytical categories, not necessarily the substantive priorities of GPT 5.5. In particular, the small Gender and Women and Schools of Law and Theology groups limit reliable comparison.

This report concerns GPT 5.5 alone. Although the keyness output compares GPT 5.5 with a "rest" corpus, the identities, sizes, token totals, and preprocessing equivalence of that comparison corpus are not provided. Accordingly, the results support a profile of this segment but do not support comparative claims about GPT 5.5 versus Gemini, Grok, Claude, or DeepSeek.

Research themes

The most dependable thematic signal is a connected ethical-juridical vocabulary rather than seven clearly separable LDA topics. The synthesis identifies a broad field organized around **ruling, principle, community, justice, protection, interpretation, and practice**. Distinctive terms within this field include **apostasy, conscience, belief, woman, humanistic_sharia, interpretation, and reform**.

The LDA export is not sufficiently reliable for fine-grained topic interpretation. Although the run specifies seven topics, the payload contains duplicated topic identifiers and labels, 15 topic records, 26 topics in coverage, inconsistent prevalence values, and no reported C_v, C_NPMI, U_Mass, perplexity, or stability metrics. It also flags low coherence for topics at indexes 2 and 4 and substantial overlap among several topic pairs. The topic distribution showing 95 dominant chunks for one topic, 41 for another, and 35 for a third can be reported descriptively, but it should not be treated as evidence of seven distinct and stable themes.

Within these limitations, the corpus supports five issue-specific thematic readings.

Gender and women's rights

Gender discourse is visible but comparatively underrepresented. The relevant vocabulary includes **woman**, **gender**, **gender justice**, **divorce**, **property**, **inheritance**, **patriarchal**, and **women's rights**. The most directly associated topic-level keywords include woman, justice, scholar, right, reform, jurisprudence, authority, dignity, principle, and gender.

The corpus includes only 6 chunks classified as Gender and Women, so this issue cannot be weighted against the other four issues with confidence. Its average sentiment score is 0.1364, with all 6 chunks labeled positive, but this is too small a base for substantive conclusions. The evidence indicates that GPT 5.5 frames women's rights through justice, dignity, reform, and interpretation; it does not establish the model's position on particular gender-related rulings.

Religious rights

Religious rights are the most visibly developed issue in the segment. The dominant terms include **apostasy**, **religious freedom**, **freedom of conscience**, **belief**, **coercion**, **minority**, **protection**, **dhimmi**, and **citizenship**. Apostasy occurs 419 times, while protection occurs 913 times, religious 851 times, minority 532 times, conscience 279 times, and freedom 236 times.

The associated resource vocabulary differentiates several aspects of the issue. Resources on ridda emphasize apostasy, repentance, political loyalty, coercion, and protection; resources on classical religious freedom emphasize apostasy, conscience, classical jurists, and premodern context; and resources on religious diversity emphasize pluralism, 'adl, individual rights, and minority experience. This suggests that the GPT 5.5 discourse does not reduce religious rights to a single freedom claim, but repeatedly places belief and conscience in relation to communal order, political authority, and protection.

Interfaith relationships

Interfaith relationship discourse is represented primarily through the language of **minority protection**, **non-Muslim communities**, **religious diversity**, **citizenship**, **community**, and **equal citizenship**. The n-grams *non-Muslim communities* and *majority and minority* occur 106 and 89 times respectively, while *religious diversity* occurs 75 times. The co-occurrence of **equal_citizenship** and **protection** occurs 53 times, and **dhimma** and **protection** occur 50 times.

These associations indicate an interpretive framing in which interfaith relations are discussed through institutional protection, civic status, and communal coexistence. The resource-level terms **minority**, **minority protection**, **pluralism**, **individual rights**, and **minority experience** reinforce this reading. However, the data do not show whether these concepts are presented as GPT 5.5's own normative commitments, as summaries of scholarly positions, or as critical comparisons among positions.

Human dignity

Human dignity is a central bridge concept across the ethical and juridical vocabulary. The term **dignity** occurs 593 times, while the n-grams *human dignity* and *human right* occur 331 and 803 times respectively. In the co-occurrence analysis, **dignity-karāmah** has an NPMI of 0.324 across 48 observations, and **dignity-justice** appears 138 times with an NPMI of 0.2328.

The dignity discourse is also linked to **protection, justice, human rights, minority, woman, and moral responsibility**. The strongest keyness signal for **moral_responsibility** is an overrepresentation in the GPT 5.5 group, with a log-ratio of 3.57 and G^2 of 264.39, while **public_welfare** has a log-ratio of 2.49 and G^2 of 207.16. These results support an interpretation of the segment as frequently organizing humanistic discussion around ethical responsibility and welfare, but they do not by themselves demonstrate that the generated discourse consistently protects karāmah in substantive legal conclusions.

Interpretive juristic reasoning

Interpretive reasoning is structured through **law, ruling, interpretation, jurisprudence, scholar, jurist, historical context, Qur'an, theological, authority, reform, and disagreement**. Law is the most frequent term at 1,265 occurrences, followed by ethical at 1,188, interpretation at 673, ruling at 619, scholar at 533, and jurisprudence at 371.

The n-grams *historical context, quranic foundation, maqāṣid al-sharī'a, juristic preference, and human interpretation* occur 187, 130, 113, 72, and 83 times respectively. The co-occurrence of **law-state** occurs 197 times, while **justice-Qur'an** occurs 52 times and **theological-theological_school** occurs 74 times. Together, these patterns indicate that interpretation is commonly represented as contextual, institutionally situated, and ethically evaluated rather than as a purely textual retrieval process.

The keyness profile further emphasizes **ethical, moral responsibility, public welfare, disagreement, authority, citizenship, and worship**. Conversely, **verse, dhimmi, quranic_foundations_sharia, advocate, integrate, and adapt** are relatively less represented in the reported comparison. Because the comparison corpus is unspecified, these contrasts should be treated as descriptive lexical differences rather than definitive model-specific tendencies.

Key concepts from TF-IDF and frequency analysis

The highest TF-IDF terms are **protection** with a weight of 15.935, **community** at 13.2028, **apostasy** at 13.1923, **interpretation** at 11.9976, **dignity** at 11.4039, **justice** at 10.7794, **right** at 10.5533, **ruling** at 10.5176, **minority** at 10.3555, and **scholar** at 10.016. The frequency ranking similarly places law, ethical, protection, religious, community, interpretation, ruling, dignity, justice, and historical at the top.

The most informative multi-word concepts include *human right, human dignity, historical context, family law, universal declaration, ethical responsibility, religious freedom, quranic foundation, freedom of conscience, maqāṣid al-sharī'a, gender justice, due process, and protection of life*. These phrases connect rights language with Islamic ethical and juridical vocabularies rather than presenting them as wholly separate domains.

Several relational patterns are especially important for a humanistic interpretation:

- **Dignity and karāmah** connect universal human worth with an Islamic ethical term.
- **Justice and mercy** associate legal evaluation with moral restraint and compassion.
- **Coercion and conscience** place religious freedom in tension with compulsion and authority.
- **Community and individual** frame rights through the relationship between collective order and personal agency.
- **Law and state** connect juristic reasoning to political institutionalization.
- **Right and woman** provide a direct lexical link between rights discourse and gender concerns.
- **Protection and violence** position protection as a response to harm, although the data do not establish how consistently that relationship is normatively resolved.

The keyness results show relatively greater use of **ethical, treat, moral responsibility, protection, public welfare, community, mean, coercion, disagreement, authority, citizenship, and worship** in the reported GPT 5.5 group. Terms relatively less represented include **align, allow, verse, dhimmi, era, integrate, advocate, lens, quranic foundations of sharia, provide, ensure, adapt, and reveal**. The TF-IDF and word-frequency sections reproduce the same keyness statistics, so they should be understood as one shared keyness result rather than independent confirmation.

Author and institution mapping

No usable author, organization, regional, or entity landscape is provided. The entity output contains no dominant entity types, named people, regions, or organizations, and the synthesis identifies no central resources. Consequently, the evidence cannot support claims about which scholars, institutions, regions, or intellectual networks shape the segment.

The available lexical evidence does identify recurring role terms, including **scholar, jurist, thinker, authority, and theological school**. These terms indicate that GPT 5.5 represents interpretive authority and scholarly mediation as important parts of sharia discourse, but they do not identify specific authors or map institutional influence.

Madhhab comparison is also absent. Although the corpus contains the compound token **hanafi_maliki_shafi'i_hanbali** and the co-occurrence pair **theological-theological_school**, there are no school-level frequencies, keyness statistics, topic distributions, or co-occurrence comparisons. The evidence therefore supports recognition of juristic plurality as a stated vocabulary, not a verified comparison of Hanafi, Maliki, Shafi'i, Hanbali, Ja'fari, or other traditions.

Paper clustering

The resource-level signals suggest several recurring resource groupings, but the supplied clustering output does not establish valid paper clusters. The only reported macro-cluster places all seven topic IDs into a single "Ruling / Principle / Community" cluster with prevalence 1, while a second cluster has prevalence 0 and no representative resources. This is tautological rather than evidence of meaningful separation.

A cautious resource-level synthesis can nevertheless distinguish the following overlapping areas:

- **Ridda and religious conversion**, marked by apostasy, repentance, political loyalty, coercion, and protection.
- **Religious diversity and minority status**, marked by pluralism, 'adl, individual rights, minority experience, citizenship, and protection.
- **Blasphemy and expression**, marked by blasphemy, insult, *sabb al-nabi*, speech, punishment, accusation, and freedom of expression.
- **Gender justice and women's rights**, marked by woman, reform, divorce, property, inheritance, gender justice, and patriarchal context.
- **Human rights and dignity**, marked by human rights, dignity, ethical resources, public welfare, and the modern nation-state.
- **Interpretive method and intellectual networks**, marked by scholars, thinkers, historical context, authority, interpretation, and contextualization.

These areas overlap substantially around law, protection, justice, community, dignity, and interpretation. They should therefore be treated as thematic concentrations for literature-review synthesis, not as statistically validated clusters of mutually exclusive papers.

Affective framing

The overall sentiment score is mildly positive at 0.1576, with 173 of 200 chunks labeled positive, 18 neutral, and 9 negative. This pattern describes lexical affect in academic prose; it does not demonstrate that GPT 5.5 adopts a substantively humanistic, rights-affirming, or ethically balanced position.

The clearest contrast occurs in the blasphemy and expression topic. Its average sentiment score is -0.0217, with 7 negative, 10 neutral, and 2 positive chunks out of 19. This concentration indicates discourse tension around insult, punishment, speech, authority, offense, and freedom of expression. It should be interpreted as an affective or rhetorical tension signal, not as evidence for a particular legal ruling.

Gender and Women has an average score of 0.1364 across only 6 chunks, while Religious Rights has a lower average of 0.102 across 30 chunks and includes 10 neutral and 1 negative chunk. The small size of the gender category and the absence of sentence-level validation details make these subgroup comparisons indicative rather than conclusive.

Research gaps and analytical limitations

The principal substantive gap is the absence of a balanced five-issue comparison. Religious rights are supported by extensive vocabulary and resource-level evidence, whereas gender and women's rights have only 6 classified chunks. Interfaith relationships and human dignity are present through minority protection, citizenship, dignity, *karāmah*, justice, and human-rights terminology, but they are not represented by separately validated issue-level measures.

Juristic plurality remains underexamined. The analysis does not compare madhhabs or theological schools, does not normalize school names and transliteration variants, and does not test whether terms associated with particular traditions occur in different ethical or juridical contexts. This prevents a strong conclusion about whether GPT 5.5 preserves *ikhtilāf* or merely mentions plurality.

The classification distribution cannot be treated as substantive emphasis because cross-validated macro-F1, per-class precision and recall, confusion matrices, and fold variability are absent. The dominance of Foundations of Islamic Law at 123 chunks may reflect taxonomy design or corpus composition rather than a validated model tendency.

Several quantitative outputs require technical verification before stronger interpretation:

- The LDA records contain inconsistent identifiers, prevalence values, topic counts, and quality information.
- NPMI values above 1 are reported, including 1.3315 for *north and west*, despite the stated theoretical range of -1 to 1.
- The text network has density 0.978, identical degree centrality of 1 across listed nodes, and identical betweenness scores of 0.00054, which makes hub and bridge interpretations unreliable.
- Multiple-testing procedures, normalized frequencies, confidence intervals, and false-discovery-rate adjustments are not reported.
- The payload was truncated, with chart datasets and display tables removed, limiting auditability of fine-grained results.
- The analysis does not distinguish model-generated claims from quoted, cited, or paraphrased scholarly positions in the underlying materials.

These limitations do not erase the observed lexical patterns, but they constrain the level of inference. In particular, positive sentiment, high-frequency ethical vocabulary, and co-occurrence patterns should not be equated with normative validity or with successful humanistic reinterpretation.

Synthesis

The GPT 5.5 segment represents sharia through a strongly interconnected ethical-juridical lexicon. Law, ethical reasoning, protection, community, interpretation, dignity, justice, minority status, and ruling form the principal vocabulary through which the five requested humanistic issues are organized.

Religious rights and interfaith relationships are the most extensively evidenced areas, especially through apostasy, conscience, coercion, minority protection, citizenship, pluralism, and non-Muslim communities. Human dignity operates as a cross-cutting normative concept, linked directly to karāmah, justice, human rights, protection, and moral responsibility. Gender and women's rights appear through woman, reform, gender justice, family law, property, inheritance, and dignity, but the six-chunk support base requires substantial caution.

Interpretive juristic reasoning is represented as contextual and ethically mediated. Historical context, maqāṣid al-sharīʿa, juristic preference, scholars, authority, disagreement, public welfare, and moral responsibility suggest an orientation toward negotiation among text, history, institutional authority, and human consequences. The associated co-occurrences--particularly dignity-karāmah, justice-mercy, coercion-conscience, community-individual, and law-state--provide the strongest evidence for how humanistic and juridical concepts are relationally assembled.

The profile is therefore best characterized as an ethically framed, protection-oriented, and community-conscious representation of sharia that repeatedly engages rights, dignity, conscience, and interpretive disagreement. It does not yet establish that GPT 5.5 resolves tensions among gender equality, religious freedom, interfaith citizenship, human dignity, and juristic plurality in a consistent or normatively defensible way. A stronger literature-review conclusion requires corrected topic and n-gram outputs, validated classification, explicit madhhab-sensitive comparison, representative excerpt audits, and matched cross-model evidence.

3.4. Gemini 3.1 Pro preview

Purpose: Academic Literature Review Synthesizer

Corpus overview

The analysis covers 156 scoped chunks from 100 resources, totaling 139,434 words for Gemini 3.1 Pro preview. The run applies topic modeling, TF-IDF, word frequency, sentiment, n-gram, co-occurrence, classification, network, and phrase-chunking analyses. Because the active scope contains only Gemini and excludes a global baseline, the results describe Gemini's corpus profile and its supplied aggregate keyness comparison, rather than a complete model-by-model comparison.

The five requested humanistic issues are unevenly represented in the available classification output. Religious rights accounts for 30 chunks, Gender and Women for 9, Foundations of Islamic Law for 43, Reform and Legal ethics for 49, and Schools of Law & Theology for 25. Interfaith relations, human dignity, and interpretive juristic reasoning are not provided as separate validated classes; their coverage must therefore be inferred cautiously from keywords, topics, phrases, and co-occurrences.

Research themes

The strongest recurring orientation is a historically situated and jurisprudential treatment of sharia. Gemini's most prominent vocabulary includes *law* with 859 occurrences, *historical* with 604, *jurisprudence* with 467, *scholar* with

420, *ruling* with 283, and *interpretation* with 244. TF-IDF and keyness reinforce this pattern: relative to the supplied remainder corpus, Gemini significantly overuses *law* ($G^2 = 144.76$; log-ratio = 0.72), *historical* ($G^2 = 131.46$; log-ratio = 0.87), and *jurisprudence* ($G^2 = 121.65$; log-ratio = 0.97).

A second theme concerns religious liberty, apostasy, conscience, and the authority of the state. *Apostasy* appears 304 times and is a high-weight TF-IDF term, while topic-distinctive vocabulary includes *religious_freedom*, *belief*, *personal_conscience*, *state*, and *human_right*. Co-occurrences connect *apostasy* with *death_penalty* and *treason*, indicating that the discourse places doctrinal questions alongside penal, political, and historical framings rather than treating apostasy solely as an abstract matter of belief.

A third theme addresses minority protection and interfaith relations. Relevant vocabulary includes *minority* (177 occurrences), *pluralism* (170), *diversity* (168), *religious_freedom* (148), *protection* (259), and *interfaith_relation* among the distinctive topic terms. Representative resources address non-Muslim worship and property, historical Muslim and non-Muslim interactions, ethical treatment of non-Muslims, minority protections, and interfaith dialogue. However, the dataset does not supply a dedicated interfaith classification count, so the breadth of this theme cannot be expressed as a validated issue prevalence.

Gender and women's rights form a smaller but substantively identifiable theme. The classified Gender and Women category contains 9 chunks, and its distinctive vocabulary includes *woman*, *gender_justice*, *female_autonomy*, *inheritance*, *property*, *patriarchal*, and *modern_feminist*. The available evidence suggests attention to legal autonomy, financial independence, inheritance, and the historical framing of gender norms, but the small category size makes prevalence and sentiment estimates comparatively uncertain.

Human dignity functions as a cross-cutting ethical concept rather than a separately measured domain. It appears 160 times, occurs in topic keywords and distinctive terms, and forms part of the vocabulary of resources concerning universal respect, interfaith relations, minority protection, and human rights. Its repeated association with *right*, *protection*, *justice*, *human*, and *ethical* supports interpreting dignity as a mediating concept linking juridical analysis to humanistic evaluation, although the data do not establish that Gemini consistently prioritizes dignity in every issue area.

Interpretive juristic reasoning is one of the clearest structural features of the corpus, despite lacking a separate classification category. The network's highest weighted nodes are *law* (2,030), *historical* (1,408), *jurisprudence* (1,101), *scholar* (936), *ruling* (714), and *interpretation* (673). Gemini also overuses *quranic_foundations_sharia* ($G^2 = 124.86$; log-ratio = 2.39), *universal_declaration_human_rights* ($G^2 = 119.66$; log-ratio = 1.35), and the combined madhhab reference *hanafi_maliki_shafi'i_hanbali* ($G^2 = 85.17$; log-ratio = 1.82). This combination indicates a discourse that repeatedly places scriptural foundations, historical context, human-rights instruments, and madhhab plurality into the same interpretive frame.

Key concepts from TF-IDF and phrase analysis

The TF-IDF profile is dominated by *law*, *historical*, *jurisprudence*, *apostasy*, *scholar*, *right*, *state*, *protection*, *ruling*, *theological*, *humanistic*, *ethical*, *pluralism*, *interpretation*, and *dignity*. This lexical pattern presents sharia as a field of legal and ethical interpretation shaped by historical development, scholarly authority, state power, and rights discourse.

Frequent multi-word expressions further connect Islamic legal reasoning to human-rights and maqāṣid-oriented language. Prominent n-grams include *human right* (783 occurrences), *historical context* (231), *quran foundation* (225), *human dignity* (159), *religious freedom* (132), *maqāṣid al-sharia* (122), *gender justice* (80), *interfaith relation* (62),

and *freedom of belief* (53). These phrases support an interpretation of Gemini's humanistic reframing as an attempt to mediate between inherited legal categories and contemporary ethical vocabularies.

Several co-occurrences identify more specific conceptual relationships. *Justice 'adl* occurs with *mercy raḥma* at an NPMI of 0.5571, *family law* with *reform* at 0.4185, and *maqāṣid al-sharia* with *objective law* at 0.5887. The pairing of *offender* and *repent* (0.555), together with *capital punishment* and *repent* (0.4786), indicates that penal questions are linguistically connected to repentance and rehabilitation-related framing. These associations are evidence of lexical proximity, not proof of a settled normative position.

The supplied n-gram output requires technical caution. Some values exceed the stated NPMI range of -1 to +1, including *ash'arī and maturidi* at 1.2373 and *international human right* at 1.1364. These values should be treated as unvalidated association outputs until the distinction between frequency, significance, and NPMI is corrected.

Author and institution mapping

No reliable author, institution, regional, or organizational entity mapping is present in the analytics payload. The corpus does identify scholarly roles and traditions through terms such as *scholar*, *jurist*, *thinker*, *theological school*, *classical fiqh*, and the Hanafi, Maliki, Shafi'i, and Hanbali references. These are discourse markers rather than verified author or institutional metadata, so they should not be used to infer representation, influence, or scholarly balance.

The resource-level vocabulary points to several intellectual settings: modern underexplored sharia theorists, historical religious tolerance, universal respect in Islam, state enforcement of *sabb al-nabi* regulations, interfaith dialogue, religious freedom, and sectarian minority protection. The titles and distinctive terms indicate topical coverage, but they do not establish source quality, citation validity, or balanced representation of classical and contemporary scholarship.

Paper clustering

The topic model does not currently provide a dependable basis for treating the corpus as a set of distinct paper clusters. Although the run specifies seven topics, the payload contains repeated topic identifiers and labels, reports 15 topic entries in one section and 26 topics in coverage, and duplicates prevalence values across multiple keyword representations. The synthesis output also assigns all seven topics to one macro-cluster with prevalence 1 and assigns a second cluster prevalence 0, indicating a degenerate clustering stage rather than meaningful paper groupings.

The dominant-topic distribution nevertheless reveals concentration in a limited set of recurring discursive areas. The largest reported topic representation covers 64 chunks, or 41.03%, followed by 42 chunks, or 26.92%; the remaining reported topics account for 32.05% collectively. The largest topic representation is associated with law, historical context, scholars, jurisprudence, dignity, protection, and humanistic reform, but its reported coherence is low at 0.296 and it overlaps substantially with other topics.

Accordingly, the most defensible exploratory groupings are thematic rather than formal paper clusters: apostasy and conscience; blasphemy and freedom of expression; gender and legal autonomy; minority and interfaith protection; and interpretive pluralism across legal and theological schools. Topic overlap is especially pronounced around protection, community, ruling, justice, dignity, and ethical principles, suggesting that these themes are interconnected rather than cleanly separable.

Comparative and sentiment signals

Gemini's supplied keyness profile distinguishes its language from the aggregate remainder in several ways. It overuses *regard* ($G^2 = 197.47$; log-ratio = 2.61), *advocate* (155.31; 2.66), *traditional* (137.98; 1.93), *complex* (129.96; 2.26), *orientalist_perspective* (100.94; 1.70), *rigid* (92.40; 1.93), and *humanistic* (78.18; 1.22). It underuses *community* (166.74; -1.59), *conscience* (67.40; -1.88), *authority* (95.24; -2.53), *worship* (82.18; -4.44), and *communal* (65.69; -2.37). These contrasts suggest a relatively legal-historical and advocacy-oriented framing, with less frequent communal, worship-centered, and authority-centered wording than the comparison corpus.

Sentiment is predominantly positive in lexical terms: the mean VADER compound score is 0.2877, with 142 positive, 10 neutral, and 4 negative chunks. This should not be interpreted as evidence that the discourse is ethically affirmative, humane, or rights-supporting, because VADER measures lexical and sentence-level polarity rather than juridical validity or moral endorsement.

The class-level distribution shows a meaningful difference in framing intensity. Schools of Law & Theology has the highest mean score at 0.4196, followed by Reform and Legal ethics at 0.3289 and Gender and Women at 0.2804. Religious rights is notably lower at 0.132, with 18 positive, 9 neutral, and 3 negative chunks, while the blasphemy and *sabb al-nabi* topic has a near-neutral mean of 0.0438. These lower values may reflect the contentious vocabulary of coercion, punishment, offense, apostasy, and state enforcement rather than a validated negative stance toward religious rights.

Research gaps

The principal analytical gap is the absence of a validated five-issue evidence matrix. Gender and women's rights have only 9 classified chunks, while religious rights has 30; interfaith relations and human dignity are embedded in broader themes without dedicated counts, and interpretive juristic reasoning is visible through vocabulary but not separately quantified. A follow-up analysis should apply consistent issue-specific rules, manually audit representative passages, and report issue prevalence, keyness, topic association, and sentiment together.

The classifier's reliability cannot be assessed because macro-F1, precision, recall, fold variability, confidence scores, and error samples are absent. The Gender and Women category should therefore be treated as low-confidence until its nine chunks are manually validated, and all class distributions should be reported as descriptive rather than definitive.

Topic interpretation also requires reconstruction. The selected K, topic identifiers, prevalence fields, and coverage counts are inconsistent; C_v, C_NPMI, U_Mass, log-perplexity, and seeded stability are not reported. The low coherence of 0.296 and extensive topic overlap mean that the current latent-topic results are exploratory and should not be presented as stable literature clusters.

Further gaps concern comparability and preprocessing. No direct results are supplied for the other named models, no normalized cross-model issue comparison is available, and unequal chunk or resource sizes may affect frequencies and topic weights. Malformed compound tokens, including variants of madhhab names and human-rights phrases, may also distort lexical and network results. Network density is very high at 0.8808, so centrality should not be treated as substantive conceptual importance without threshold and sensitivity analyses.

Synthesis

Gemini's representation of sharia is best characterized as historically conscious, jurisprudentially oriented, and explicitly mediated through human-rights and ethical vocabularies. The discourse repeatedly connects law and

rulings with scholars, interpretation, historical context, pluralism, protection, dignity, maqāṣid, and madhhab references. Its treatment of the five humanistic concerns is uneven: religious rights and interpretive reasoning are strongly visible, gender is present but comparatively small, while interfaith relations and dignity operate mainly as cross-cutting themes rather than separately measured domains.

The most substantive humanistic tensions appear around apostasy, conscience, state authority, blasphemy, freedom of expression, minority protection, and gendered legal autonomy. Co-occurrences such as justice-mercy, family law-reform, apostasy-death penalty, and offender-repentance indicate that Gemini often frames sharia through balancing concepts rather than isolated rules. This is consistent with an interpretive posture that places inherited jurisprudence in dialogue with maqāṣid, historical context, universal human-rights language, and ethical protection.

These findings should remain qualified. The supplied data support a strong descriptive account of Gemini's vocabulary and framing, but they do not establish the model's doctrinal accuracy, its consistency across prompts, or its superiority to other models. A defensible next stage is to rebuild the topic schema, validate preprocessing and n-gram scoring, audit the five issue categories, and apply the same normalized measures across every model before drawing comparative conclusions about AI-mediated humanistic reinterpretations of sharia.

3.5. Grok 4.3

Purpose: Academic Literature Review Synthesizer

Corpus overview

The Grok 4.3 segment contains 151 chunks drawn from 100 resources and 127,524 words. Preprocessing enabled stopword removal and lemmatization, with a maximum vocabulary of 2,000 features. The analysis combined topic modeling, TF-IDF, word frequency, sentiment, n-gram, co-occurrence, classification, network, and phrase-based methods.

The corpus is segmented by model, and the displayed segment contains Grok 4.3 only. Accordingly, most findings describe Grok's internal lexical and thematic profile. The keyness results provide a contrast with a rest-of-corpus baseline, but the peer-model composition, sample sizes, prompt equivalence, and resource balance are not documented; comparative claims should therefore be treated as indicative rather than definitive.

Research themes

The most prominent vocabulary concerns interpretation, ruling, justice, community, principle, historical context, protection, religious rights, ethics, dignity, minorities, and pluralism. Word frequency analysis places interpretation first with 485 occurrences, followed by ruling with 426, justice with 417, community with 411, principle with 393, and historical with 366. TF-IDF similarly prioritizes justice, interpretation, ruling, community, principle, historical, protection, and apostasy.

The synthesis suggests a discourse organized around interpretive mediation: legal rulings are repeatedly related to ethical principles, community practice, protection, and contextual adaptation. This is supported by the network's weighted hubs, where interpretation, community, justice, principle, ruling, allow, historical, religious, protection, and ethical occupy the leading positions. However, the network has a density of 0.9559 and an average clustering coefficient of 0.9588, while several degree scores equal 1; these construction characteristics limit the value of degree and betweenness as discriminators.

The LDA output should not be treated as a validated account of seven distinct topics. Although the run parameters state a topic count of 7, the payload reports 15 topic records, a topic count of 26 in coverage, repeated topic identifiers and labels, and duplicated prevalence values. Quality metrics are empty, no seeded stability score is available, and several topics are flagged as low-coherence or overlapping, including coherence values of 0.2797 and 0.276. The broad cluster labeled "Community / Ruling / Principle," with a reported prevalence of 89.81%, is therefore best understood as a provisional lexical grouping rather than a distinct substantive discovery.

Within that limitation, the topic evidence points to two broad areas of emphasis:

- **Community, protection, and ethical interpretation.** Recurrent terms include community, ruling, principle, justice, practice, minority, religious, pluralism, woman, conscience, and protection.
- **Belief freedom and contested religious offences.** Apostasy, belief, conscience, punishment, coercion, conversion, blasphemy, and freedom of belief form a more identifiable but narrower cluster. Apostasy appears in 15 of 151 dominant-topic chunks, or 9.93%, and is repeatedly associated with punishment and belief.

Key concepts from TF-IDF and frequency analysis

Grok's corpus-wide lexical profile is strongly process-oriented. The most distinctive comparative terms are integrate, adapt, allow, focus, address, maintain, contrast, ensure, justice, external, guidance, principle, broad, respect, and practice. Integrate has a log-likelihood G^2 of 317.61 and a log-ratio of 3.82 against the rest-of-corpus baseline; adapt has G^2 316.80 and log-ratio 2.73. These results support the interpretation that Grok frequently frames sharia-related reasoning through integration, adaptation, contrast, and practical guidance.

This process orientation should be distinguished from corpus-wide prominence. Justice, interpretation, and ruling are highly frequent and important within Grok's segment, but the supplied evidence does not establish that they are uniquely Grok-specific. By contrast, the keyness results show that Grok underuses law, state, classical_fiqh, classical_jurist, maqāsid, quranic, dhimma, citizenship, human_right, and woman relative to the comparison corpus. Woman occurs 46 times in the Grok segment compared with 684 occurrences in the rest-of-corpus count, with a log-ratio of -1.68.

The n-gram output identifies recurring phrases such as human right, historical context, Quran foundation, minority and majority, universal declaration, human dignity, principle of justice, freedom of belief, minority communities, interpretation traditions, and textual fidelity. The reported field labeled significance includes values above 1, including 1.3233, despite the methodology describing NPMI as bounded between -1 and 1. These values should not be interpreted as valid NPMI until the implementation or field labeling is audited.

Evidence across the five humanistic issues

Gender and women's rights

Gender and women's rights are present but substantially underrepresented. The Gender and Women classification contains only 2 of 151 chunks, and the term woman appears 46 times in Grok's segment compared with 684 in the rest-of-corpus keyness counts. The topic vocabulary includes woman, gender, justice, reform, authority, dignity, property, inheritance, protection, patriarchal, and historical context, indicating that the issue is framed through rights, reform, and jurisprudential interpretation.

The evidence does not support a strong conclusion about Grok's substantive orientation toward women's rights. The small two-chunk sample is not comparable with the Foundations of Islamic Law class, which contains 58 chunks, or the Reform and Legal Ethics class, which contains 52. Gender findings should therefore be treated as

exploratory, and additional balanced coverage would be required before characterizing a stable model-level position.

Religious rights

Religious rights are more visible than the gender issue. The Religious Rights class contains 24 chunks, and the most prominent associated terms include apostasy, belief, conscience, coercion, punishment, religious, conversion, and freedom of belief. Apostasy is the leading term in the narrower belief-related topic, and the co-occurrence pair apostasy-punishment appears 21 times.

The lexical pattern suggests that Grok represents religious rights primarily through the tension between belief and coercive or punitive legal responses. This interpretation is reinforced by the topic vocabulary surrounding conscience, context, political authority, individual rights, and religious freedom. It should not be read as evidence that the model resolves these tensions in a particular theological or juridical direction; the quantitative data show recurrent problem framing rather than normative correctness.

Interfaith relations

Interfaith relations appear in selected resources and resource-level distinctive terms, but they are not among the strongest corpus-wide lexical signals. Representative resources include "Strategies for fostering interfaith dialogue under sharia," "Practical steps toward interfaith understanding in sharia settings," and "Contrasting sharia approaches to religious pluralism worldwide." Resource-level terms include interfaith, mutual recognition, religious, participant, feedback, evaluation, and prohibition.

The distinction between resource framing and corpus-wide prominence is important. Interfaith is not among the leading global terms or top co-occurrence pairs, whereas majority-minority occurs 19 times and individual_right-pluralism occurs 20 times. The evidence therefore supports an interpretation of interfaith relations as a contextual application of broader pluralism, minority protection, and community ethics, rather than as the dominant lexical center of the segment.

Human dignity

Human dignity is a substantive ethical connector across the corpus. The term dignity occurs 222 times, and human dignity is a recurring n-gram with 56 occurrences. Dignity also co-occurs with karāmah 19 times, while dignity-principle occurs 51 times. Related terms include justice, protection, minority, human, ethical, right, and welfare.

These relationships suggest that Grok's humanistic reinterpretation links dignity to protection and justice rather than presenting it as an isolated abstract value. The resource vocabulary further includes karāmah, inherent worth, dignified treatment, and dignity-centered approaches. This supports a cautious interpretation of dignity as a mediating concept between legal reasoning, minority rights, and ethical treatment, although the data do not establish how consistently the model prioritizes karāmah when it conflicts with inherited rulings.

Interpretive juristic reasoning

Interpretive juristic reasoning is the most consistently visible of the five issues. Interpretation is the most frequent term at 485 occurrences and the leading weighted network node, while ruling occurs 426 times. Jurisprudence,

scholar, historical, context, text, principle, Quran, ethical, adapt, and guidance also recur across frequency, TF-IDF, topic, and co-occurrence outputs.

Grok's keyness profile indicates that this reasoning is often expressed as a process of integration, adaptation, contrast, maintenance, and guidance. The model appears, on the available evidence, to frame juristic reasoning less through explicit classical institutional vocabulary and more through contextual and practical mediation. This interpretation is supported by the underrepresentation of `classical_fiqh`, `classical_jurist`, `maqāṣid`, `quranic`, and named madhhab combinations, including the combined Hanafi-Maliki-Shafi'i-Hanbali term. It should not be taken to mean that Grok rejects classical jurisprudence; rather, the segment gives less lexical prominence to explicit historical-doctrinal labels than to interpretive operations.

Paper clustering and classification

The taxonomy assigns 58 chunks to Foundations of Islamic Law, 52 to Reform and Legal Ethics, 24 to Religious Rights, 12 to Schools of Law and Theology, and 2 to Gender and Women. These categories provide a useful descriptive map of corpus coverage, with foundations and reform together accounting for 110 of 151 chunks.

The category distribution should not be confused with validated semantic clustering. The supplied output reports class counts but does not provide cross-validated macro-F1, per-class precision and recall, a confusion matrix, or a low-confidence flag. The reliability of the classifier is therefore unknown, and the Gender and Women category is too small to support meaningful predictive evaluation.

Topic overlap further limits claims about discrete paper or resource groups. The reported overlap scores reach 0.5789 between topic groupings involving dignity, rights, minorities, authority, and justice, and 0.55 between groupings involving protection, community, religious practice, and dignity. These overlaps indicate that the corpus's five issues are interwoven rather than cleanly separable, but the duplicated LDA records prevent firm conclusions about independent clusters.

Author and institution mapping

No reliable author, institution, organization, regional, or named-jurist landscape is available in the supplied analytics. The entity outputs are empty, and the analysis does not provide a source-level author or institutional metadata map. Consequently, no evidence-based claim can be made about which scholars, madhhabs, institutions, or geopolitical contexts most influence Grok's discourse.

The vocabulary does contain references to scholar, jurist, thinker, Hanafi, Ash'ari, Maturidi, and intellectual networks. These terms demonstrate that juristic and theological authority is represented, but they do not identify the actual authors or establish the relative influence of particular schools. A future review should standardize transliteration and separately extract madhhab, jurist, author, institution, and region entities.

Research gaps

The strongest empirical gap is gender coverage. Only 2 chunks are classified as Gender and Women, and woman is underrepresented in the comparative keyness results. The current corpus cannot support a robust account of how Grok handles women's rights, gender justice, inheritance, authority, family law, or patriarchal interpretation.

A second gap concerns direct comparative evaluation. The segment is scoped to Grok 4.3, while the baseline behind keyness is not fully documented. A valid model comparison would require aligned prompts, equal or

normalized resource weights, comparable chunking, five-issue coverage measures, peer sample sizes, topic distributions, and classifier quality metrics.

A third gap is methodological auditability. LDA quality metrics and stability are missing; topic identifiers and prevalence values are inconsistent; n-gram significance appears incompatible with the stated NPMI range; and the high-density network provides limited structural discrimination. The 100-resource and 151-chunk relationship is also unexplained, leaving open the possibility that some resources contribute disproportionately to frequencies, topics, and sentiment.

Finally, sentiment should not be used as evidence of humanistic or ethical approval. VADER labels 149 of 151 chunks positive, with an average score of 0.2854, but this may reflect academic evaluative language and recommendation verbs. The result measures lexical polarity, not theological balance, juristic legitimacy, pluralism, or commitment to human dignity.

Synthesis

Grok 4.3 represents sharia primarily through an interpretive and process-oriented vocabulary linking ruling, principle, justice, community, protection, historical context, and adaptation. Its most distinctive comparative signals are integrate, adapt, contrast, maintain, ensure, guidance, and respect, while explicit references to classical jurisprudence, state authority, maqāṣid, named jurists, citizenship, and women are comparatively less prominent.

Across the five humanistic issues, religious rights and interpretive reasoning are the strongest observable domains. Apostasy, belief, conscience, punishment, and coercion form a recurrent rights-related problem space, while interpretation, ruling, jurisprudence, scholar, historical context, and adaptation organize the model's approach to juristic reasoning. Human dignity functions as an ethical bridge through dignity-karāmah, human dignity, justice, protection, and minority-rights associations.

Interfaith relations are visible mainly through selected resources and pluralism-related relations rather than dominant corpus-wide terminology. Gender and women's rights are the least secure area of interpretation because of the two-chunk classification support and the comparative underuse of woman. Overall, the evidence supports describing Grok's discourse as adaptive, integrative, and ethically framed, but not as comprehensively humanistic across all five issues.

A defensible literature-review conclusion is therefore provisional: Grok frequently mediates between legal principles and contemporary ethical concerns, yet its humanistic reinterpretation remains unevenly evidenced. Stronger conclusions require corrected topic serialization, validated classification, balanced gender coverage, documented peer baselines, resource-level normalization, standardized Islamic legal terminology, and explicit comparison across all model groups.

4. CORPUS OVERVIEW

Composition and size of the corpus in the current scope: document types, chunk length statistics and coverage of the project's coding parameters. Generated October 3, 2026 10:30.

4.1. Document Composition

Table 1. Document composition by type in the current scope.

DOCUMENTS BY TYPE	DOCUMENTS
Article Import	500
Total	500

4.2. Chunk Length (Words)

Table 2. Chunk length statistics in words for the current scope.

CHUNKS	TOTAL WORDS	MEAN	MEDIAN	MIN	MAX
761	734,447	965.1	1140.0	38	1,481

Word counts are whitespace-based estimates (space count + 1) computed on chunk contents.

4.3. Parameter Coverage — Model

Table 3. Parameter coverage — Model: chunks per value in the current scope.

MODEL	CHUNKS
GPT 5.5	200
Gemini 3.1 Pro preview	156
DeepSeek V4 Pro	154
Grok 4.3	151
Claude Opus 4.7	100

Top 5 values by chunk count within the current scope.

5. CONCLUSION AND RECOMMENDATIONS

Conclusion

Across models, AI-generated discourse consistently constructs sharia as a negotiated field linking law, rights, ethics, historical context, and communal order rather than as a single fixed code. The most stable semantic core is the connection among apostasy, religious freedom, conscience, minority protection, dignity, and human rights. "Human right" is a highly significant n-gram across all groups, with frequencies ranging from 304 for Grok 4.3 to 803 for GPT 5.5, while apostasy is prominent in TF-IDF and frequency results for every model. This convergence indicates that reinterpretation is organized around the tension between inherited rulings and contemporary claims concerning personal belief, equality, expression, and state power.

The models nevertheless differ in emphasis and interpretive posture. Claude Opus 4.7 concentrates strongly on pluralism and minority rights: its leading LDA topic accounts for 73.49% of topic prevalence, and its classification is led by Reform and Legal ethics at 43% and Religious rights at 20%. Its vocabulary combines apostasy, woman, dignity, misconception, reform, and scholar, suggesting a corrective and humanistic framing grounded in juristic debate. DeepSeek V4 Pro is more legally and textually intensive: "law" is its most frequent term at 486 occurrences, while its topics distribute more evenly across Sharia, dignity, pluralism at 49.67%, punishments and human rights at 26.07%, and apostasy and religious liberty at 24.25%. It therefore foregrounds institutional law, punishment, Quranic reasoning, maqāṣid, and the protection of life, intellect, lineage, and property.

GPT 5.5 produces the most expansive and internally connected discourse. Its network density is 0.9780 and average clustering is 0.9798, substantially above Claude's 0.8522 density and 0.8719 clustering. Its lexical profile is dominated by law at 1,169 occurrences, ethical at 1,029, protection at 734, and religious at 624, while noun-phrase extraction places sharia at 1,138 and justice at 556. This model integrates foundational law, rights, conscience, dignity, community, reform, and gender justice, but its LDA distribution is comparatively concentrated: Religious Minority Rights Under Sharia accounts for 43.13%, with State Institutions and Islamic Rights at 23.37%.

Gemini 3.1 Pro preview is the most explicitly comparative and school-conscious. Its classification gives substantial weight to Reform and Legal ethics at 31.4% and Schools of Law & Theology at 16%, while its strongest keyword signals include Hanafi, Maliki, Hanbali, Shafi'i, Ash'ari, Maturidi, and Mu'tazila. Its n-grams connect maqāṣid, objective law, justice and mercy, family-law reform, repentance, and pluralism. Gemini therefore offers the clearest basis for preserving ikhtilāf, although its high positive sentiment score of 91.0% may indicate a tendency toward harmonizing language rather than sustained representation of disagreement.

Grok 4.3 emphasizes interpretation, justice, scholarly networks, and implementation. Interpretation is its most frequent term at 350, justice at 326, and its classification is divided between Foundations of Islamic Law at 39.2% and Reform and Legal ethics at 35.1%. Its distinctive co-occurrences include individual right-pluralism, majority-minority, dignity-karāmah, and justice-mercy. This makes Grok particularly useful for practical, context-sensitive and evaluative analysis, though its 98.7% positive sentiment suggests possible under-detection of conflict and legal risk.

A central limitation is methodological: LDA reports substantial topic overlap and low coherence, especially where dignity, protection, punishment, and rights recur across themes. Very positive sentiment across all models should not be interpreted as ethical agreement or juridical safety. The corpus also appears highly standardized across models, so frequency differences partly reflect corpus size and repeated source structures rather than wholly distinct reasoning styles.

Recommendations

- **High -- Build a juristic plurality and disagreement protocol.** Require every generated interpretation to distinguish Quranic principles, hadith-based arguments, madhhab positions, contemporary reformist readings, and state practice. The protocol should explicitly identify areas of agreement, disagreement, evidentiary uncertainty, and unresolved ikhtilāf rather than compressing them into a single "humanistic" conclusion.
- **High -- Make maqāṣid and karāmah operational rather than rhetorical.** Use a structured evaluation sequence that tests protection of life, religion, intellect, family, property, dignity, conscience, and equal civic standing. The recurrent links among dignity-karāmah, protection, justice, and human rights support this framework, but outputs should show how each principle affects a concrete legal or ethical judgment.
- **High -- Separate descriptive fiqh from normative recommendation.** For apostasy, blasphemy, gender justice, and minority status, require distinct sections for classical doctrine, school-based variation, historical context, modern legislation, and proposed reinterpretation. This is especially important because apostasy and punishment dominate keyness across models, while topic overlap risks conflating legal categories and contemporary ethical preferences.
- **Medium -- Use model specialization through ensemble review.** Use Gemini for madhhab and theological comparison, Claude for dignity-centered pluralism, DeepSeek for legal-textual and maqāṣid analysis, Grok for contextual implementation and scholarly-network mapping, and GPT 5.5 for broad synthesis. A final adjudication layer should reconcile their outputs and flag substantive divergence rather than selecting the most fluent answer.
- **Medium -- Calibrate sentiment and safety assessment.** Positive sentiment ranges from 77.0% for Claude to 98.7% for Grok, which may reflect respectful prose rather than low controversy. Add human review and targeted checks for coercion, incitement, dehumanization, unjustified certainty, and inadvertent endorsement of state punishment.
- **Medium -- Improve the quantitative pipeline.** Consolidate overlapping LDA topics, report corpus-normalized frequencies, and evaluate robustness across alternative topic counts and preprocessing choices. Network density and clustering are extremely high for all models, so future analysis should emphasize edge strength, bridge terms, and community structure rather than density alone.
- **Low -- Expand underrepresented perspectives.** Increase coverage of women scholars, Shi'i and minority madhāhib, regional traditions, lived minority experiences, and non-state legal practice. This would reduce the present concentration on apostasy, punishment, and human-rights comparison while strengthening the project's commitment to epistemic humility and genuine plurality.